

Fall 2013

B-TREES

(LOOSELY BASED ON THE COW BOOK: CH. 10)

Motivation

Consider the following table:

```
CREATE TABLE Tweets (  
    uniqueMsgID INTEGER,      -- unique message id  
    tstamp      TIMESTAMP,    -- when was the tweet posted  
    uid         INTEGER,      -- unique id of the user  
    msg         VARCHAR (140), -- the actual message  
    zip         INTEGER       -- zipcode when posted  
);
```

Consider the following query, Q1: `SELECT * FROM Tweets
WHERE uid = 145;`

And, the following query, Q2: `SELECT * FROM Tweets
WHERE zip BETWEEN 53000 AND 54999`

Ways to evaluate the queries, efficiently?

1. Store the table as a heapfile, scan the file. I/O Cost?
2. Store the table as a **sorted file**, binary search the file. I/O Cost?
3. Store the table as a heapfile, build an **index**, and search using the index.
4. Store the table in an **index** file. The entire tuple is stored in the index!

Index

- Two main types of indices
 - **Hash** index: good for equality search (e.g. Q1)
 - **B-tree** index: good for both range search (e.g. Q2) and equality search (e.g. Q1)
 - Generally a hash index is faster than a B-tree index for equality search
- Hash indices aim to get $O(1)$ I/O and CPU performance for search and insert
- B-Trees have $O(\log_F N)$ I/O and CPU cost for search, insert and delete.

What is in the index


- Two things: **index key** and **some value**
 - Insert(indexKey, value)
 - Search (indexKey) -> value (s)
- What is the index key for Q1 and Q2?
- Consider Q3:

```
SELECT * FROM Tweets  
WHERE uid = 145 AND  
zip BETWEEN 53000 AND 54999
```

- Value:
 - Record id
 - List of record id
 - The entire tuple!

(Ubiquitous) B+ Tree

- Height-balanced (dynamic) tree structure
- Insert/delete at $\log_F N$ cost (F = fanout, N = # leaf pages)
- Minimum 50% occupancy (except for root).
Each node contains $d \leq \underline{m} \leq 2d$ entries.
The parameter d is called the **order** of the tree.
- Supports equality and range-searches efficiently.



Index Entries


Entries in the index
(i.e. non-leaf) pages:
(search key value, pageid)

Data Entries

Entries in the leaf pages:
(search key value, recordid)

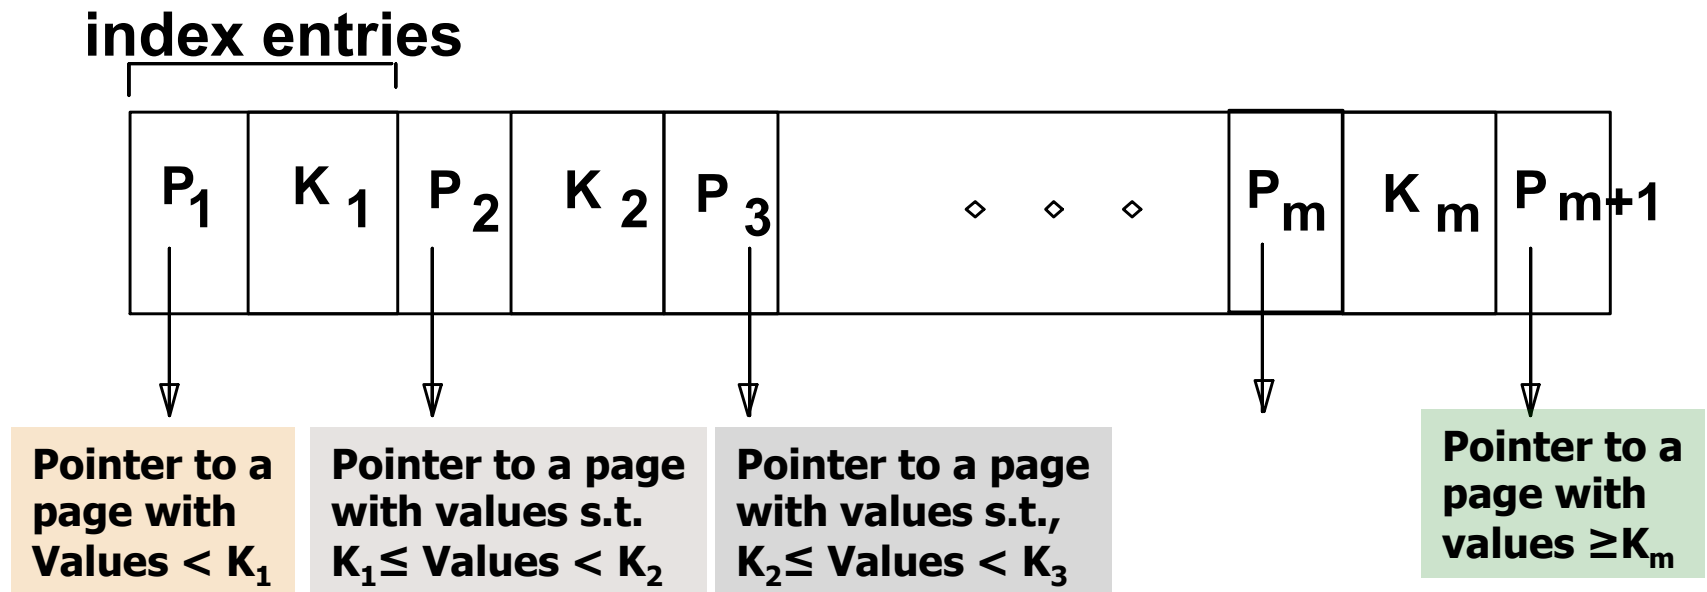
Example B+ Tree

- Search: Starting from root, examine index entries in non-leaf nodes, and traverse down the tree until a leaf node is reached
 - Non-leaf nodes can be searched using a binary or a linear search.
- Search for 5*, 15*, all data entries $\geq 24^*$




B+-tree Page Format

Non-leaf
Page



Leaf Page





B+ Trees in Practice

- Typical order: 100. Typical fill-factor: 67%.
 - average fanout = 133
- Typical capacities:
 - Height 4: $133^4 = 312,900,700$ records
 - Height 3: $133^3 = 2,352,637$ records
- Can often hold top levels in buffer pool:
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 Mbyte
 - Level 3 = 17,689 pages = 133 MBytes

B+-Tree: Inserting a Data Entry

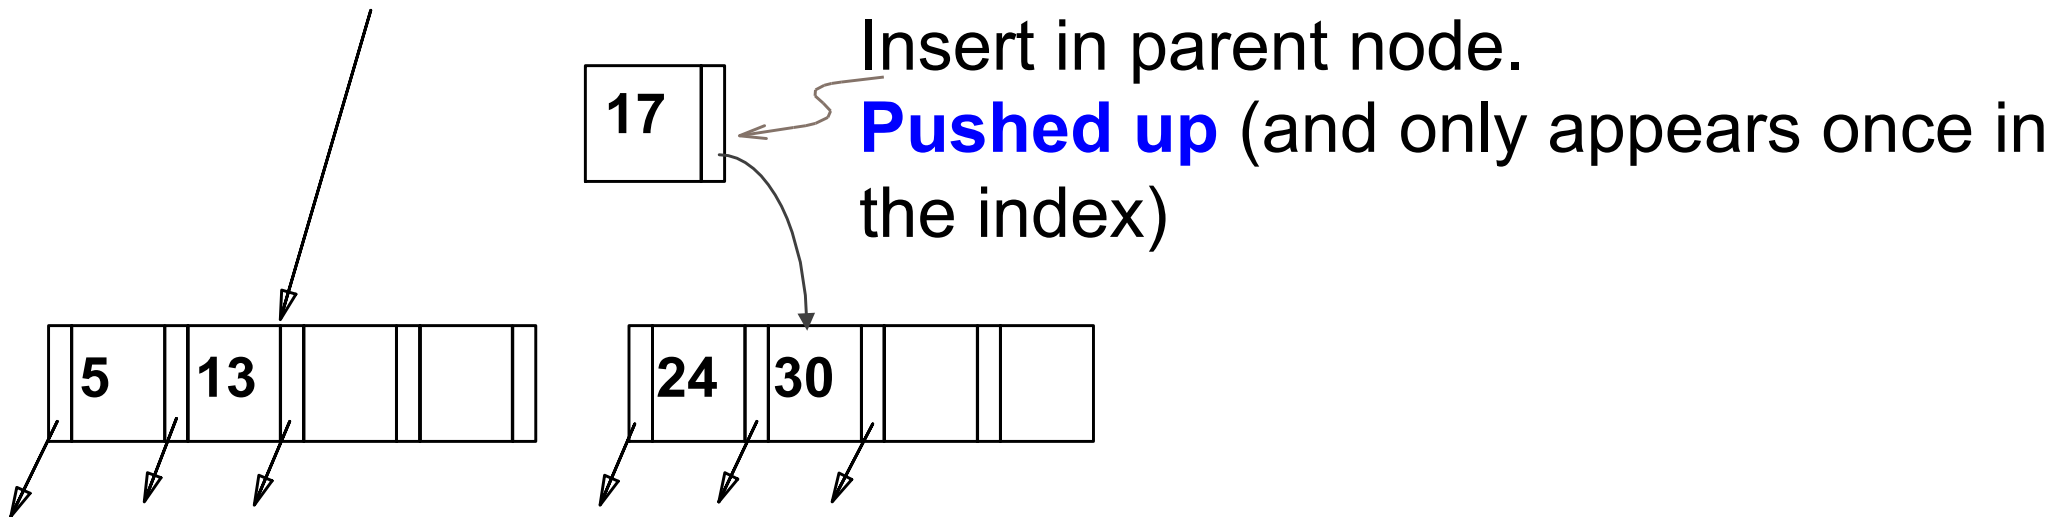
- Find correct leaf L .
- Put data entry onto L .
 - If L has enough space, *done!*
 - Else, must **split** L (into L and a new node $L2$)
 - Redistribute entries evenly, **copy up** middle key.
 - Insert index entry pointing to $L2$ into parent of L .
- This can happen recursively
 - To split non-leaf node, redistribute entries evenly, but **pushing up** the middle key. (Contrast with leaf splits.)
- Splits “grow” tree; root split increases height.
 - Tree growth: gets *wider* or *one level taller at top*.

Inserting 8* into B+ Tree




Entry to be inserted in parent node
Copied up (and continues to appear in the leaf)

Inserting 8* into B+ Tree

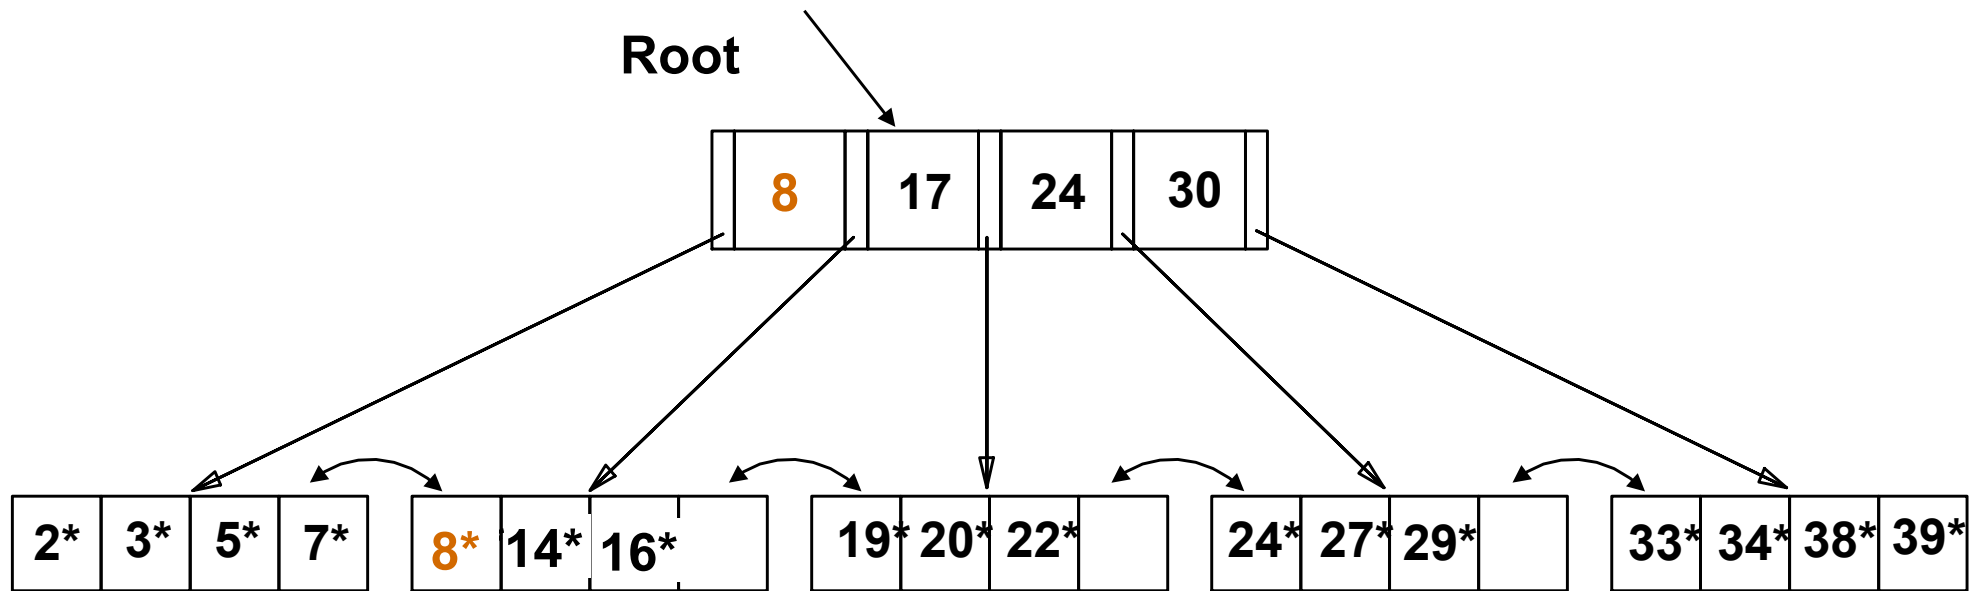


Inserting 8* into B+ Tree



- Root was split: height increases by 1
- Could avoid split by re-distributing entries with a sibling
 - Sibling: immediately to left or right, and same parent

Inserting 8* into B+ Tree




- Re-distributing entries with a **sibling**
 - Improves page occupancy
 - Usually not used for non-leaf node splits. Why?
 - Increases I/O, especially if we check both siblings
 - Better if split propagates up the tree (rare)
 - Use only for leaf level entries as we have to set pointers

B+-Tree: Deleting a Data Entry

- Start at root, find leaf L where entry belongs.
- Remove the entry.
 - If L is at least half-full, *done!*
 - If L has only **d-1** entries,
 - Try to **re-distribute**, borrowing from sibling (*adjacent node with same parent as L*).
 - If re-distribution fails, **merge** L and sibling.
- If merge occurred, must delete entry (pointing to L or sibling) from parent of L .
- Merge could **propagate** to root, decreasing height.



Deleting 22* and 20*



- Deleting 22* is easy.
- Deleting 20* is done with re-distribution.
Notice how middle key is **copied up**.


... And Then Deleting 24*

- Must merge.
- In the non-leaf node, *toss* the index entry with key value = 27




Non-leaf Re-distribution

- Tree *during deletion* of 24*.
- Can re-distribute entry from left child of root to right child.



After Re-distribution

- Rotate through the parent node
- It suffices to re-distribute index entry with key 20; For illustration 17 also re-distributed



B+-Tree Deletion

- Try redistribution with **all** siblings first, then merge. Why?
 - Good chance that redistribution is possible (large fanout!)
 - Only need to propagate changes to parent node
 - Files typically grow not shrink!

Duplicates

- Duplicate Keys: many data entries with the same key value
- Solution 1:
 - All entries with a given key value reside on a single page
 - Use overflow pages!
- Solution 2:
 - Allow duplicate key values in data entries
 - Modify search
 - Use RID to get a **unique** (composite) key!
- Use list of rids instead of a single rid in the leaf level
 - Single data entry could still span multiple pages

A Note on Order

- *Order (d)* concept replaced by physical space criterion in practice (*at least half-full*).
 - Index (i.e. non-leaf) pages can typically hold many more entries than leaf pages.
 - Leaf pages could have actual data records
 - Variable sized records and search keys mean different nodes will contain different numbers of entries.
 - Even with fixed length fields, multiple records with the same search key value (*duplicates*) can lead to variable-sized data entries (e.g. list of rids).

ISAM - Indexed Sequential Access Method

- A *static* B+-tree
 - When the index is created, build a B+-tree on the relation
 - Updates and deletes don't change the non-leaf pages.
 - Use overflow pages. Leaf pages could be empty!
- Search Cost: $\log_F N + \# \text{ overflow pages}$

