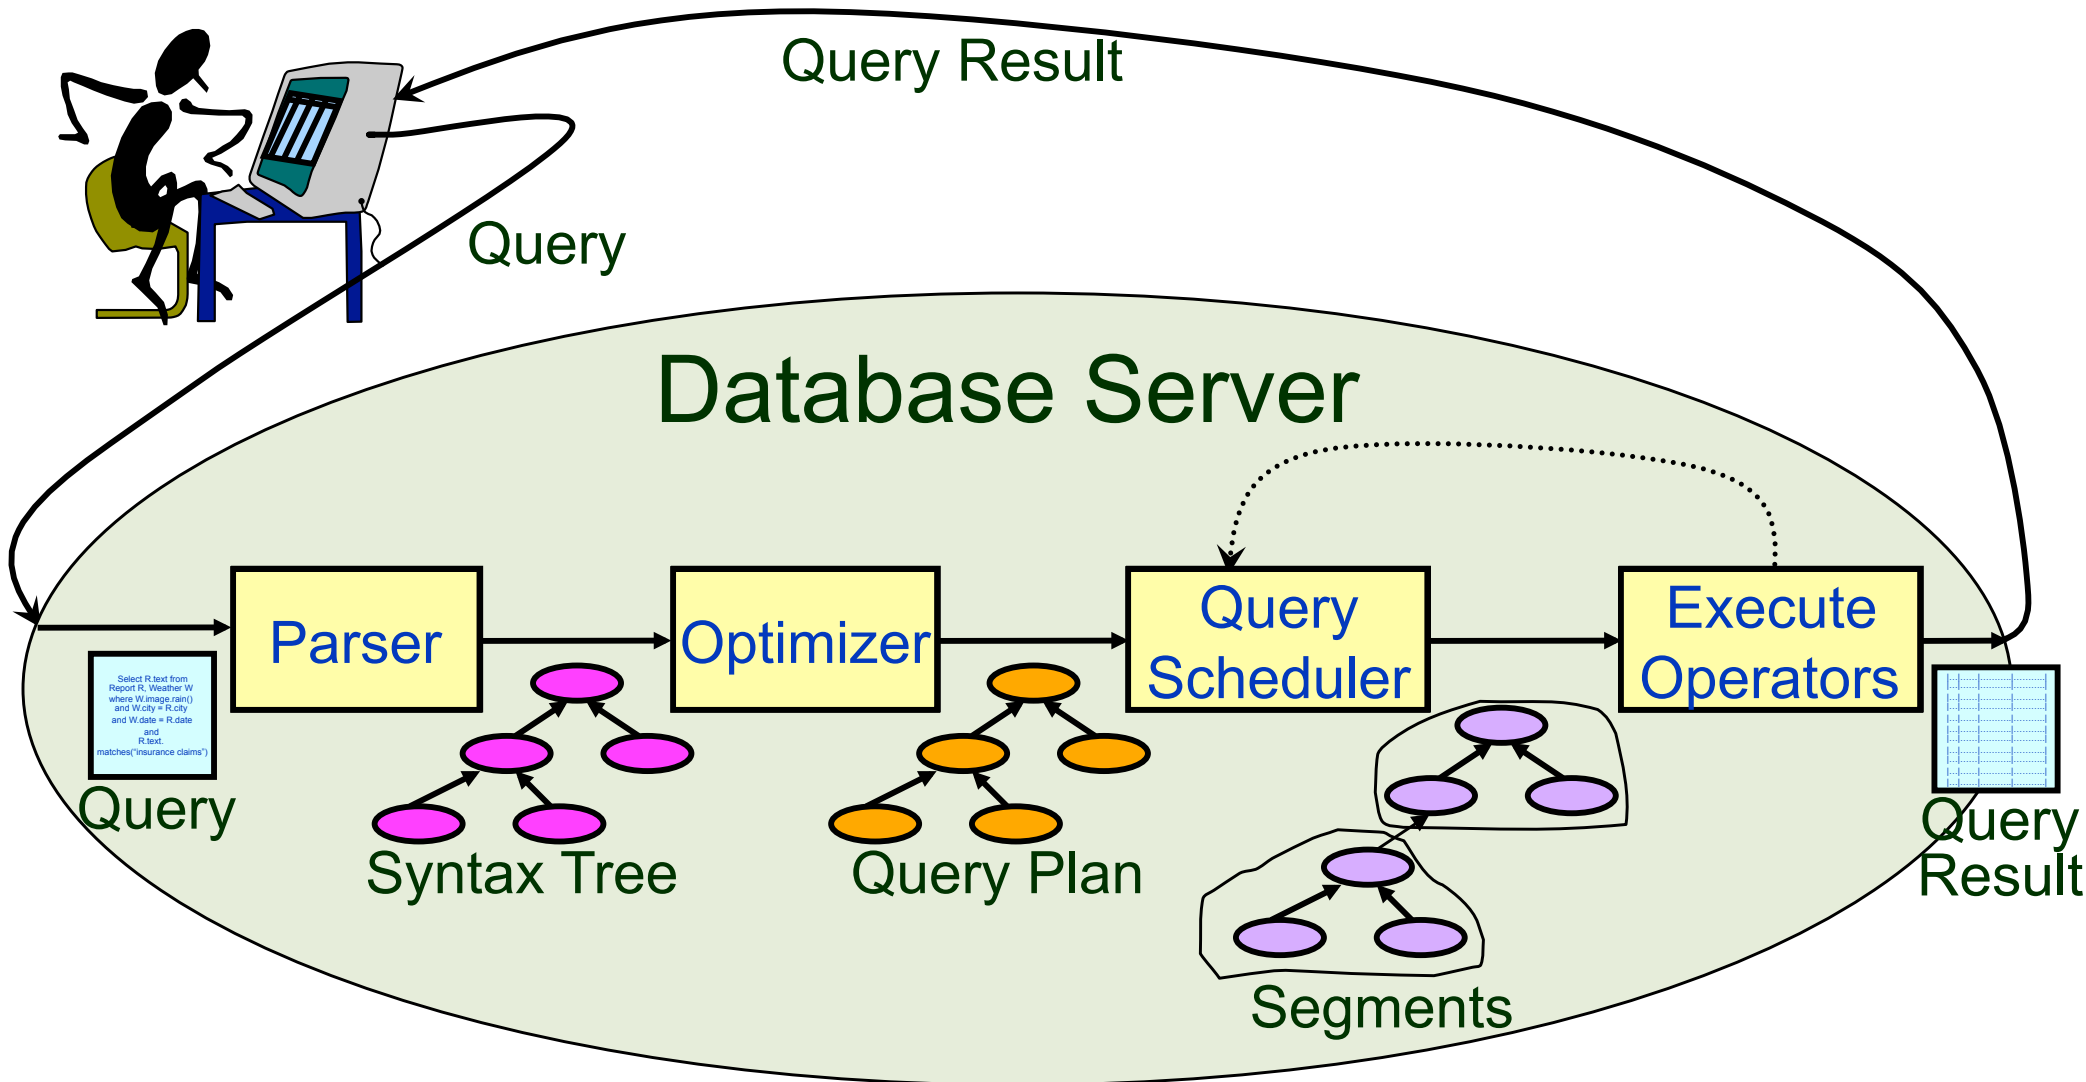


Fall 2013

# QUERY PROCESSING

**[LOOSELY BASED ON CH 12.1-12.3 AND 14 IN THE COW BOOK]**

# Life Cycle of a Query



# Problem Statement

```
CREATE TABLE User (  
    uid            INTEGER,          -- unique id of the user  
    login          VARCHAR(20)       -- unique login name  
    lname          VARCHAR(80),      -- lastname  
    fname          VARCHAR(80),      -- firstname  
    dob            DATE,             -- date of birth  
    PRIMARY KEY (uid),               -- primary key for the table  
    UNIQUE (login)                   -- twid is also unique  
);  
  
CREATE TABLE Messages (  
    uniqueMsgID    INTEGER,          -- unique message id  
    tstamp         TIMESTAMP,        -- when was the message posted  
    uid            INTEGER,          -- unique id of the user  
    msg            VARCHAR (140),    -- the actual message  
    zip            INTEGER,          -- zipcode when the message was posted  
    reposted       BOOLEAN           -- is this a reposted message?  
    PRIMARY KEY (uniqueMsgID),       -- primary key  
    FOREIGN KEY (uid) REFERENCES USER -- Foreign key to the User table  
);
```

# Problem Statement

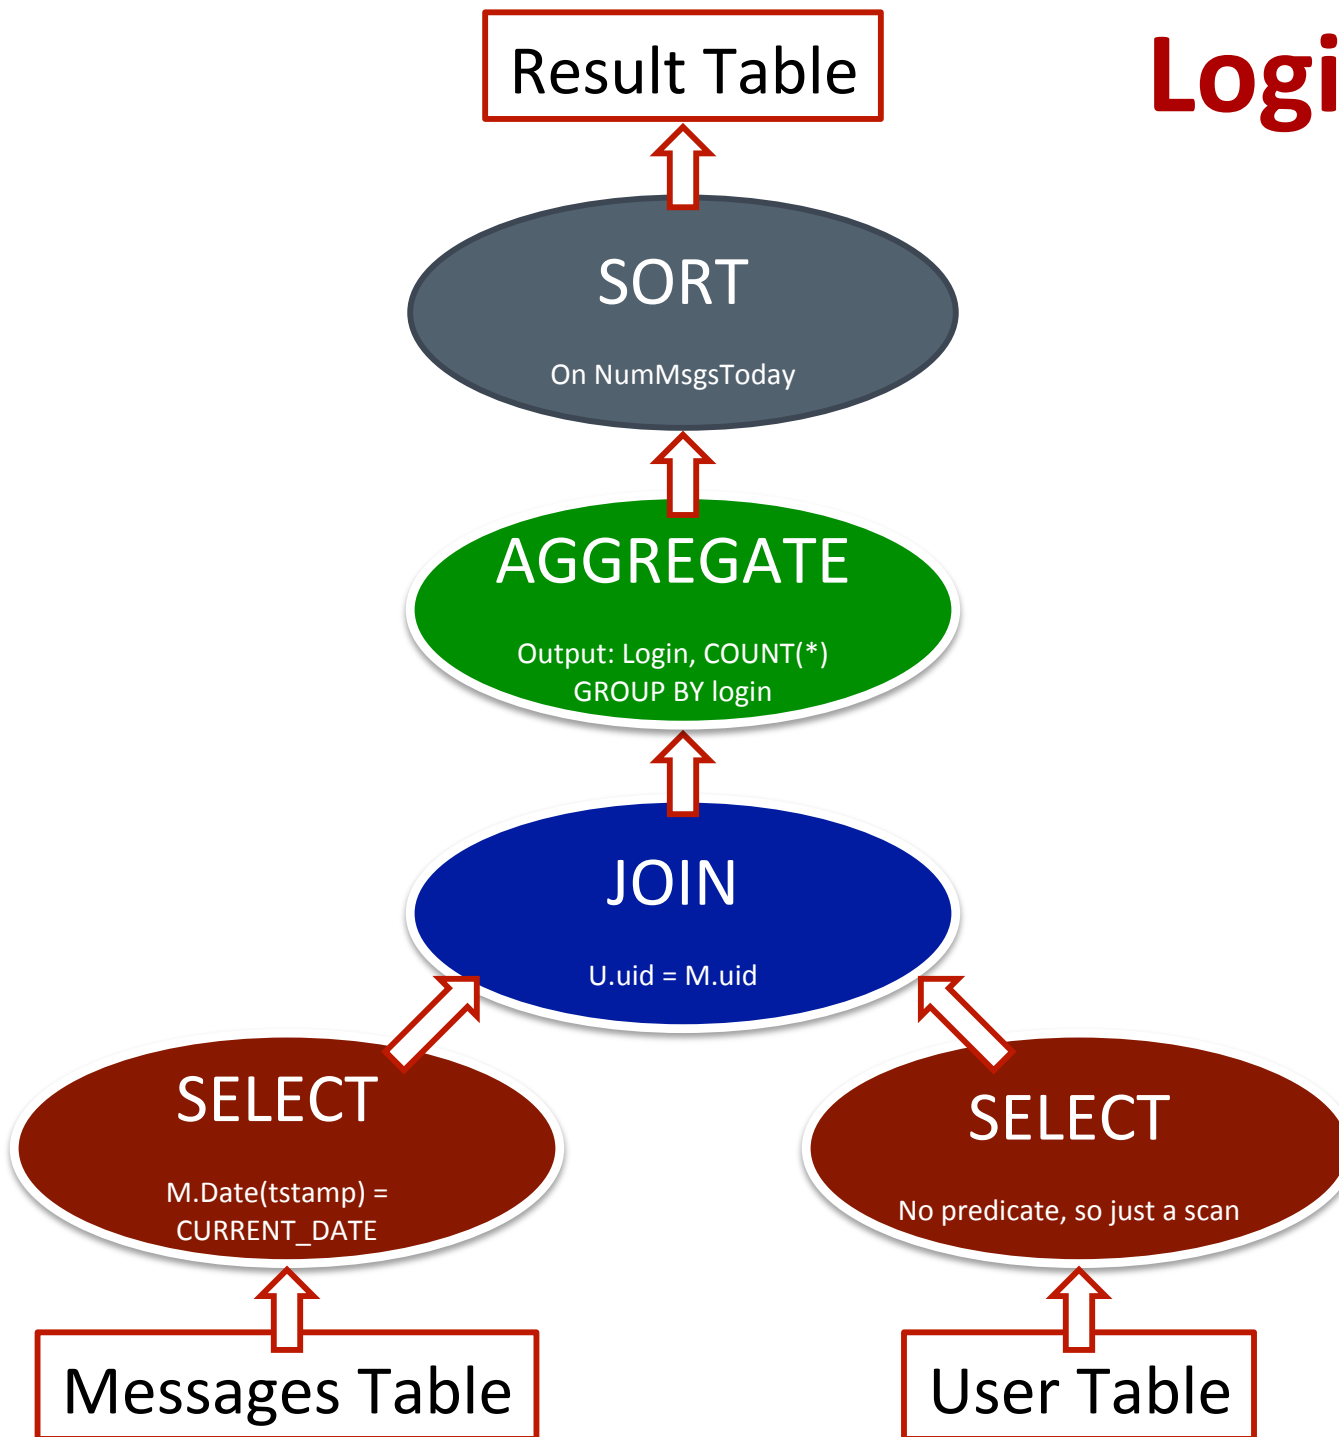
- Run the following query:

```
SELECT U.login AS login, COUNT(*) AS NumMsgsToday
FROM   User U, Messages M
WHERE  U.uid = M.uid
      AND M.Date(tstamp) = CURRENT_DATE -- select msgs posted today
GROUP BY U.login                        -- group by login
ORDER BY NumMsgsToday DESC              -- order by descending msg count
```

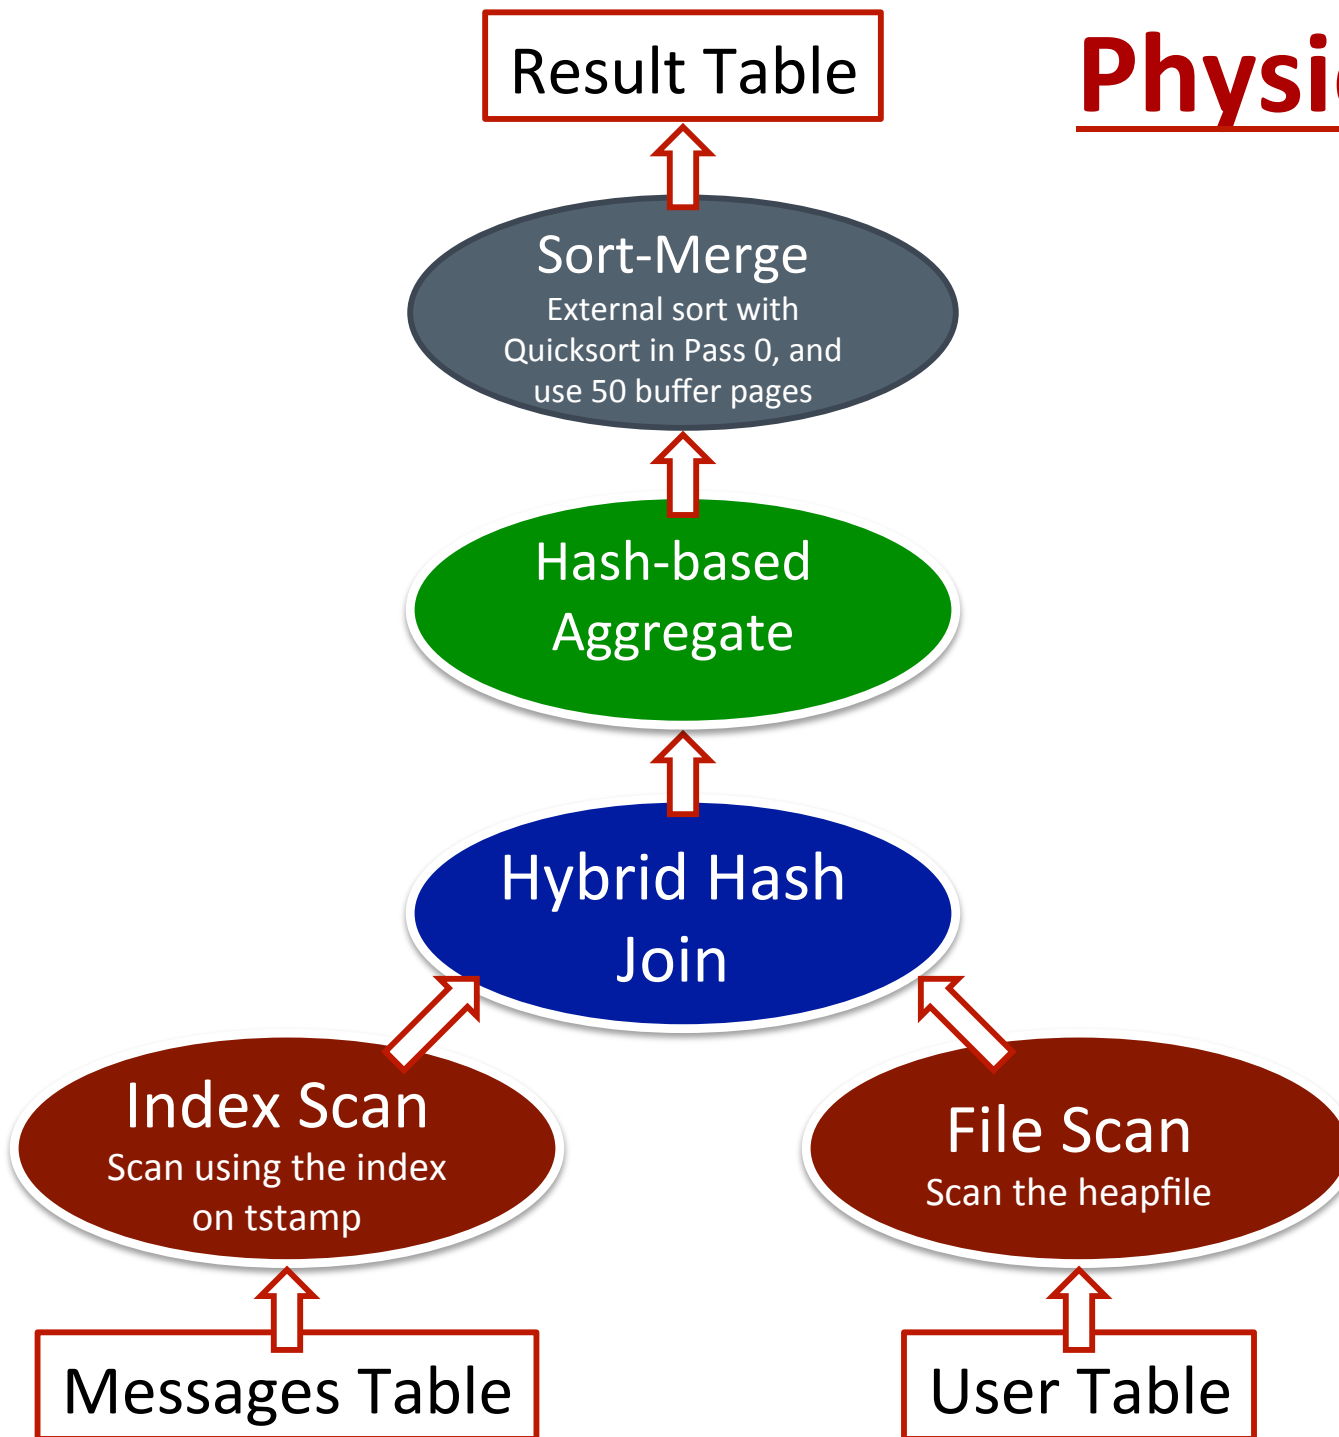
## Sample output table

login	NumMsgsToday
angelak	211
jackdr	101
petescafe	10
...	...

# Logical Query Plan



# Physical Query Plan



# Select Operation

- Algorithms: File Scan or Index Scan
- **File Scan:** Disk I/O cost:
- **Index Scan:** (on some predicate). Disk I/O cost:
  - Hash:  $O(\quad)$  can only use with equality predicates
  - B+-tree:  $O(\quad) + X$ 
    - $X$  = number of selected tuples/number of tuples per page
    - $X = 1$  per selected tuple with an unclustered index. To reduce the value of  $X$ , we could sort the rids and then fetch the tuples.
  - Bitmap Index:

# When to use a B+tree index

- Consider
  - A relation with 1M tuples
  - 100 tuples on a page
  - 500 (key, rid) pairs on a page

# data pages  
 $= 1\text{M}/100 = 10\text{K pages}$

# leaf idx pgs  
 $= 1\text{M} / (500 * 0.67)$   
 $\sim 3\text{K pages}$

	1% Selection	10% Selection
Clustered	30 + 100	300 + 1000
Non-Clustered	30 + 10,000	300 + 100,000
NC + Sort Rids	30 + ( $\sim 10,000$ )	300 + ( $\sim 10,000$ )

- ⇒ Choice of Index access plan, consider:
- 1. Index Selectivity**
  - 2. Clustering**
- ⇒ Similar consideration for hash based indices



# When can we use an index

- Notion of “index matches a predicate”
- Basically mean when can an index be used to evaluate predicates in the query

# General Selection Conditions

- Index on (R.a, R.b)
  - Hash or tree-based
- Predicate:
  - $R.a > 10$
  - $R.b < 30$
  - $R.a = 10$  and  $R.b = 30$
  - $R.a = 10$  or  $R.b = 30$
- Predicate: (p1 and p2) or p3
- Convert to Conjunctive Normal Form(CNF)  
(p1 or p3) and (p2 or p3)
- An index *matches* a predicate
  - Index can be used to evaluate the predicate

# Index Matching

- B+-tree index on  $\langle a, b, c \rangle$

- $a=5$  and  $b=3$ ?
- $a > 5$  and  $b < 3$
- $b=3$
- $a=7$  and  $b=5$  and  $c=4$   
and  **$d > 4$**
- $a=7$  and  $c=5$



**(primary conjunct)**



## Hash Idx



- Index matches (part of) a predicate

1. Conjunction of terms involving only attributes **(no disjunctions)**
2. Hash: only equality operation, predicate has all index attributes.
3. Tree: Attributes are a prefix of the search key, any ops.

# Index Matching

- A predicate could match more than 1 index
- Hash index on  $\langle a, b \rangle$  and B+tree index on  $\langle a, c \rangle$ 
  - $a=7$  and  $b=5$  and  $c=4$  Which index?
- Option1: Use either (or a file scan!)
  - Check selectivity of the primary conjunct
- Option2: Use both! Algorithm: Intersect rid sets.
  - Sort rids, retrieve rids in both sets.
  - Side-effect: tuples retrieved in the order on disk!

# Selection

- Hash index on <a> and Hash index on <b>
  - $a=7$  **or**  $b>5$  Which index?
  - Neither! File scan required for  $b>5$
- Hash index on <a> and B+-tree on <b>
  - $a=7$  **or**  $b>5$  Which index?
  - Option 1: Neither
  - Option 2: Use both! Fetch rids and union
    - Look at selectivities closely. Optimizer!
- Hash index on <a> and B+-tree on <b>
  - $(a=7$  **or**  $c>5)$  and  $b > 5$  Which index?
  - Could use B+-tree (check selectivity)

# Projection Algorithm

- Used to project the selected attributes.

**Simple case:** Example SELECT R.a, R.d.

- Algorithm: for each tuple, only output R.a, R.d

**Harder case:** DISTINCT clause

- Example: SELECT DISTINCT R.a, R.d
  - Remove attributes and eliminate duplicates
- Algorithms
  - Sorting: Sort on all the projection attributes
    - Pass 0: eliminate unwanted fields. Tuples in the sorted-runs may be smaller
    - Eliminate duplicates in the merge pass & in-memory sort
  - Hashing: Two phases
    - Partitioning
    - Duplicate elimination

# Hashing

Can  $h1 = h2$ ?

What if the hash table for a partition overflows, i.e. can't fit in memory?

$$R' = \pi_P(R)$$

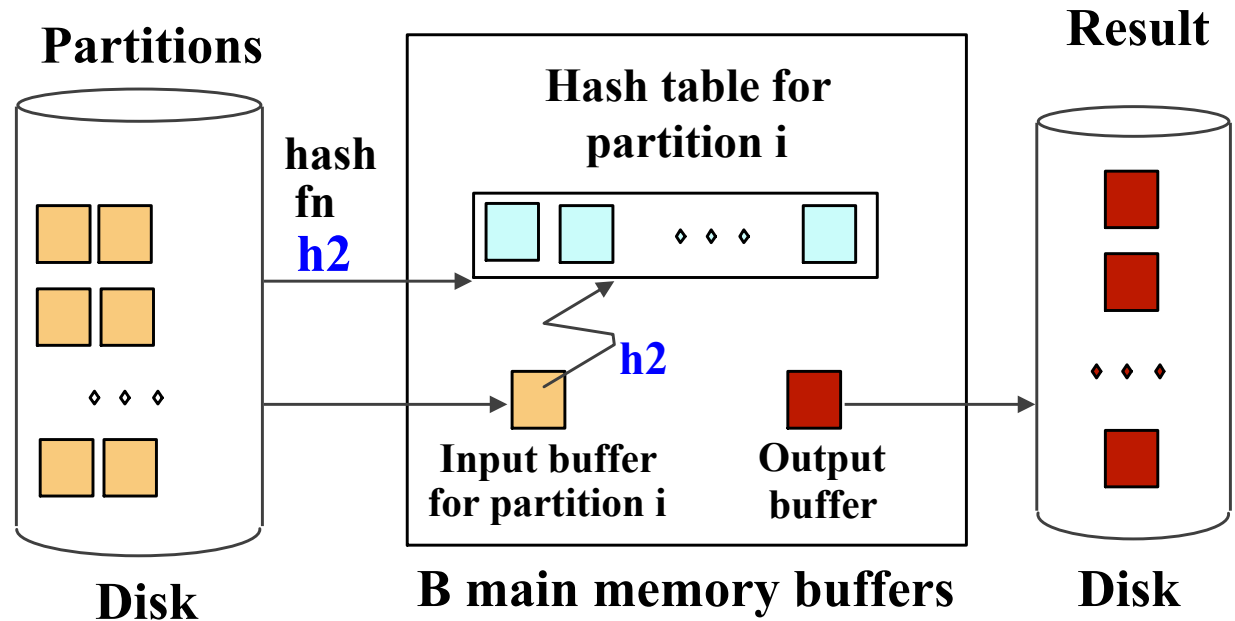

No overflows if

$$|R'| < B^2/F$$

**F = fudge factor** (to account for the hash table)

$$\text{Part. Sz, } P = |R'|/B-1$$

$$\text{Hash Tab Sz} = F * P < B$$



# Projection ...

- Sort-based approach
  - better handling of skew
  - result is sorted
  - I/O costs are comparable if  $B^2 > |R'|$
- Index-only scan
  - Projection attributes subset of index attributes
  - Apply projection techniques to data entries (much smaller!)
- If an ordered (i.e., tree) index contains all projection attributes as *prefix* of search key:
  1. Retrieve index data entries in order
  2. Discard unwanted fields
  3. Compare adjacent entries to eliminate duplicates (if required)