

SQL: Queries, Programming, Triggers

Chapter 5, Cow Book
or

See <http://sqlzoo.net/>

SQL Language

- DDL: Data definition language
- DML: Data manipulation language
- Embedded and Dynamic SQL
- Triggers
- Security
- Transaction Management
- Remote Database access

Basic SQL Query

SELECT [Optional] *target-list* **FROM** *relation-list* **WHERE** *qualification*

Attributes from input relations

List of relations

Attr1 **op** Attr2
OPS: <, >, =, <=, >=, <>
Combine using AND, OR, NOT

- Semantics/Conceptual evaluation strategy:
 - Compute the cross-product of *relation-list*.
 - Discard resulting tuples if they fail *qualifications*.
 - Delete attributes that are not in *target-list*.
 - If DISTINCT is specified, eliminate duplicate rows.
- Not an efficient evaluation plan! (Optimizer picks efficient plans)

Example of Conceptual Evaluation

```
SELECT S.name, A.hours FROM Senators S, Attendance A
WHERE S.ssn = A.ssn and A.date = '24-Sept-2010'
```

<u>ssn</u>	name	email	age	income	<u>ssn</u>	<u>date</u>	hours
11-111	Bob	bob@ca.gov	51	100.1	11-111	12-Aug-2010	1.1
22-222	Jane	jane@mi.gov	54	130.1	33-333	24-Sept-2010	4.1
33-333	Jane	jane@wi.gov	51	99.8			

(ssn)	name	email	age	income	(ssn)	date	hours
11-111	Bob	bob@ca.gov	51	100.1	11-111	12-Aug-2010	1.1
22-222	Jane	jane@mi.gov	54	130.1	11-111	12-Aug-2010	1.1
33-333	Jane	jane@wi.gov	51	99.8	11-111	12-Aug-2010	1.1
11-111	Bob	bob@ca.gov	51	100.1	33-333	24-Sept-2010	4.1
22-222	Jane	jane@mi.gov	54	130.1	33-333	24-Sept-2010	4.1
33-333	Jane	jane@wi.gov	51	99.8	33-333	24-Sept-2010	4.1

Find senators who attended the '24-Sept-2010' session

```
SELECT S.name  
FROM Senators S, Attendance A  
WHERE S.ssn = A.ssn and date = '24-Sept-2010'
```

- Add DISTINCT to this query. Effect?
- Replace S.name by S.ssn.
Effect of adding DISTINCT to this query

RA: $\pi_{\text{name}} (\sigma_{\text{date}='24-Sept-2010'}(\text{Senators} \bowtie \text{Attendance}))$

- Equivalent SQL?
- What is the schema of Senators \bowtie Attendance?

A Note on Range Variables

- Needed only if the same relation appears twice in the FROM clause.

```
SELECT S.name, A.hours  
FROM Senators S, Attendance A  
WHERE S.ssn = A.ssn and date = '24-Sept-2010'
```

OR

```
SELECT Senators.name, Attendance.hours  
FROM Senators , Attendance  
WHERE Senators.ssn = Attendance.ssn and  
Attendance.date = '24-Sept-2010'
```

*It is good style,
however, to
always use
range variables!*

Expressions and Strings

```
SELECT S.name, S.age, age1=S.age+2, S.income/S.age AS iar  
FROM Senators S  
WHERE S.sname LIKE 'Ja_%Doe'  
ORDER BY S.name
```

- Illustrates use of arithmetic expressions and string pattern matching
- `AS` and `=` are two ways to name fields in result.
- `LIKE` is used for string matching. ``_`` stands for any one character and ``%`` stands for 0 or more arbitrary characters.
- `Collation`: sort order for character sets

Find senators who attended either the '24-Sept-2010' or '25-Sept-2010' session

- UNION: Compute the union of two *union-compatible* sets of tuples
 - Same number/types of fields.
- Also available: INTERSECT and EXCEPT (What do we get if we replace UNION by EXCEPT?)
- SQL oddities: duplicates with union, except, intersect
 - Default: eliminate duplicates!
 - Use ALL to keep duplicates

```
SELECT S.ssn
FROM Senators S, Attendance A
WHERE S.ssn = A.ssn
and (A.date = '24-Sept-2010' or
     A.date = '25-Sept-2010')
```

```
SELECT S.ssn
FROM Senators S, Attendance A
WHERE S.ssn = A.ssn
and A.date = '24-Sept-2010'
UNION
SELECT S.ssn
FROM Senators S, Attendance A
WHERE S.ssn = A.ssn
and A.date = '25-Sept-2010'
```


Find senators who attended both the '24-Sept-2010' **and** '25-Sept-2010' session

- INTERSECT: Compute the intersection of any two *union-compatible* sets of tuples.
- In the SQL/92 standard, but some systems don't support it.

```
SELECT S.ssn
FROM Senators S, Attendance A1,
      Attendance A2
WHERE S.ssn = A1.ssn and S.ssn = A2.ssn
and    A1.date = '24-Sept-2010'
and    A2.date = '25-Sept-2010'
```

```
SELECT S.ssn, S.name
FROM Senators S, Attendance A
WHERE S.ssn = A.ssn
and    A.date = '24-Sept-2010'
```

INTERSECT

```
SELECT S.ssn, S.name
FROM Senators S, Attendance A
WHERE S.ssn = A.ssn
and    A.date = '25-Sept-2010'
```

Key field!

What happens if S.name is used

Nested Queries

Sailors (sid, sname, rating, age)

Reserves (sid, bid, day)

Boats (bid, bname, color)

Find names of sailors who've reserved boat #103

Can you rewrite
this to not use
a nested query?

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN (SELECT R.sid
                FROM Reserves R
                WHERE R.bid=103)
```

- Powerful feature of SQL
 - WHERE clause can itself contain an SQL query!
 - Actually, so can FROM and HAVING clauses
- To find sailors who've *not* reserved #103, use NOT IN
- Conceptual Evaluation: *nested loops* For each Sailors tuple, check the qualification by computing the subquery.

Nested Queries with Correlation

~~---Find names of sailors who've reserved boat #103---~~

Find names of sailors with exactly one reservation for boat #103

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS (SELECT UNIQUE R.bid
              FROM Reserves R
              WHERE R.bid=103 AND S.sid=R.sid)
```

Why R.bid?

- EXISTS tests if the set is not empty
- UNIQUE returns true if the row appears only once
- Illustrates why, in general, subquery must be re-computed for each Sailors tuple.

More on Set-Comparison Operators

- We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS and NOT UNIQUE.
- Also available: *op* ANY, *op* ALL, *op* is <, ≤, >, ≥, =, ≠
- Find sailors whose rating is greater than that of some sailor called Horatio:

```
SELECT *  
FROM Sailors S  
WHERE S.rating > ANY (SELECT S2.rating  
                      FROM Sailors S2  
                      WHERE S2.sname='Horatio')
```

Rewriting Except Queries Using NOT IN

Sailors (sid, sname, rating, age)
Reserves (sid, bid, day)
Boats (bid, bname, color)

Find sailors (sid) who've reserved some boat for '24-Sept-2010' but have no reservations for '09-Oct-2010'

```
SELECT S.sid
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid
      AND R.day='24-Sept-2010'
EXCEPT
(SELECT S2.sid
FROM   Sailors S2, Reserves R2
WHERE  S2.sid=R2.sid
      AND R2.day='09-Oct-2010')
```

```
SELECT S.sid
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid
      AND R.day='24-Sept-2010'
      AND S.sid NOT IN
      (SELECT S2.sid
FROM   Sailors S2, Reserves R2
WHERE  S2.sid=R2.sid
      AND R2.day='09-Oct-2010')
```

- Similarly, INTERSECT queries re-written using IN.

Division in SQL

Sailors (sid, sname, rating, age)
Reserves (sid, bid, day)
Boats (bid, bname, color)

Find sailors who've reserved **all** boats.

Without EXCEPT:

```
(1) SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
      ((SELECT B.bid
        FROM Boats B)
      EXCEPT
      (SELECT R.bid
        FROM Reserves R
        WHERE R.sid=S.sid))
```

Sailors S such that ...
there is no boat B without ...
a Reserves tuple showing S reserved B AND R.sid=S.sid))

```
(2) SELECT S.sname
FROM Sailors S
```

```
WHERE NOT EXISTS (SELECT B.bid
                   FROM Boats B
```

Sailors S such that ...

there is no boat B without ...

a Reserves tuple showing S reserved B AND R.sid=S.sid))

Aggregate Operators

```
SELECT COUNT (*) FROM Sailors S
```

```
SELECT COUNT (DISTINCT S.name)  
FROM Sailors S
```

```
SELECT AVG (S.age)  
FROM Sailors S  
WHERE S.rating=10
```

```
SELECT S.sname FROM Sailors S  
WHERE S.rating= (SELECT MAX(S2.rating) FROM Sailors S2)
```

```
COUNT (*)  
COUNT ( [DISTINCT] A)  
SUM ( [DISTINCT] A)  
AVG ( [DISTINCT] A)  
MAX (A) Can use Distinct  
MIN (A) Can use Distinct
```

single column

```
SELECT AVG ( DISTINCT S.age)  
FROM Sailors S  
WHERE S.rating=10
```

Find name & age of the oldest sailor(s)

- The first query is illegal! (wait for GROUP BY.)
- Q3 is allowed in the SQL/92 standard, but not supported in some systems

How many tuples in the result?

```
SELECT S.sname, MAX (S.age)
FROM Sailors S
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE S.age =
      (SELECT MAX (S2.age)
       FROM Sailors S2)
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE (SELECT MAX (S2.age)
       FROM Sailors S2)
      = S.age
```


GROUP BY and HAVING

- Apply aggregate to each of several *groups* of tuples
- Find the age of the youngest sailor *for each rating level*
 - Don't know: # rating levels, and rating values
 - Suppose we did know that rating values go from 1 to 10
 - we can write 10 queries that look like this (!):

For $i = 1, 2, \dots, 10$:

```
SELECT MIN (S.age)
FROM Sailors S
WHERE S.rating =  $i$ 
```

```
SELECT MIN (S.age), S.rating
FROM Sailors S
GROUP BY S.rating
```

Queries With GROUP BY and HAVING

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>

How many tuples
in the result?

- The **target-list** contains
 - Attribute names: must be a subset of *grouping-list*.
 - Terms with aggregate operations (e.g., MIN (*S.age*)).
- The **group-qualification**
 - Must have a single value per group

Conceptual Evaluation

- Cross-product -> discard tuples -> apply projection
-> partition into groups using the *grouping-list* attribute values
-> eliminate groups that don't satisfy the *group-qualification*
- Expressions in *group-qualification* have a single value per group!
 - In effect, an attribute in *group-qualification* that is not an argument of an aggregate op also appears in *grouping-list*. (SQL does not exploit primary key semantics here!)
- One answer tuple is generated per qualifying group.

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age  $\geq$  18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

- 2nd column of result is unnamed.
(Use AS to name it)

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

rating	age
1	33.0
7	45.0
7	35.0
8	55.5
10	35.0

rating	
7	35.0

Answer relation

**For each red boat, find the number of
reservations for this boat**

```
SELECT B.bid, COUNT (*) AS scout  
FROM   Sailors S, Boats B, Reserves R  
WHERE  S.sid=R.sid AND R.bid=B.bid AND B.color='red'  
GROUP BY B.bid
```

```
SELECT B.bid, COUNT (*) AS scout  
FROM   Sailors S, Boats B, Reserves R  
WHERE  S.sid=R.sid AND R.bid=B.bid  
GROUP BY B.bid  
HAVING B.color = 'red'
```

Would this work?
note: one color per bid

Find the age of the youngest sailor with age>18, for each rating with at least 2 sailors (of any age)

```
SELECT S.rating, MIN (S.age) AS MINAGE
FROM Sailors S
WHERE S.age > 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT (*) FROM Sailors S2
            WHERE S.rating=S2.rating)
```

- Subquery in the HAVING clause
- Compare this with the query where we considered only ratings with 2 sailors over 18!

Find ratings for which the average age is the minimum of the average age over all ratings

- Aggregate operations cannot be nested! WRONG:


```
SELECT S.rating
FROM Sailors S
WHERE AVG(S.age) =
      (SELECT MIN (AVG (S2.age)) FROM Sailors S2)
```

- Correct solution (in SQL/92):

```
SELECT Temp.rating, Temp.avgage
FROM (SELECT S.rating, AVG (S.age) AS avgage
      FROM Sailors S
      GROUP BY S.rating) AS Temp
WHERE Temp.avgage = (SELECT MIN (Temp.avgage) FROM Temp)
```

Null Values

- Represent
 - *unknown* (e.g., rating not assigned) or
 - *inapplicable* (e.g., no spouse's name)
- Complications with nulls:
 - Operators to check if value is/is not *null*.
 - Is *rating* > 8 true or false when *rating* is null?
 - Answer: Evaluate to unknown
 - What about **AND**, **OR** and **NOT** connectives?
 - Need 3-valued logic (true, false and *unknown*)
 - Not unknown = unknown
 - WHERE clause eliminates rows that **don't evaluate to true**
 - New operators (in particular, *outer joins*) possible/needed.



p	q	p AND q	p OR q
T	T	T	T
T	F	F	T
T	U	U	T
F	T	F	T
F	F	F	F
F	U	F	U
U	T	U	T
U	F	F	U
U	U	U	U

Outer Join

Sailors

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
58	rusty	10	35.0

Reserves

<u>sid</u>	bid	day
22	101	10/10/99

Select S.sid, R.bid
From Sailors S NATURAL **LEFT** OUTER JOIN Reserves R

Result

<u>sid</u>	bid
22	101
58	null

Similarly:

- **RIGHT OUTER JOIN**
- **FULL OUTER JOIN**

Embedded SQL

- Call SQL commands from a host language (e.g., C) program.
 - SQL statements can refer to *host variables* (including special variables used to return status).
 - Must include a statement to *connect* to the right database.
- SQL relations are (multi-) sets of records, with no *a priori* bound on the number of records. No such data structure in C.
 - SQL supports a mechanism called a *cursor* to handle this.

Cursors

- Can declare a cursor on a relation or query statement (which generates a relation).
- Can *open* a cursor, and repeatedly *fetch* a tuple then *move* the cursor, until all tuples have been retrieved.
 - Special clause, called **ORDER BY**, in cursor queries to control the order in which tuples are returned.
 - Fields in **ORDER BY** must also appear in **SELECT** clause.
- Can also modify/delete tuple pointed to by a cursor

Cursor that gets names of sailors who've reserved a red boat, in alphabetical order

```
EXEC SQL DECLARE sinfo CURSOR FOR  
  SELECT S.sname  
  FROM Sailors S, Boats B, Reserves R  
  WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'  
  ORDER BY S.sname
```

- Can we replace *S.sname* by *S.sid* in the ORDER BY clause!
 - Every column in the ORDER BY clause must appear in the SELECT clause

Integrity Constraints

- An IC describes conditions that every *legal instance* of a relation must satisfy.
 - Inserts/deletes/updates that violate IC's are disallowed.
- Types of IC's: Domain constraints, primary key constraints, foreign key constraints, general constraints.
- Can create new domains, Domain Constraints:
 - CREATE DOMAIN LegalRatings INTEGER DEFAULT 0
CHECK (VALUE >= 1 and VALUE <=10)
 - Create Table Sailor (... , rating LegalRatings, ...)
 - Underlying domain is still Integers for comparison
- Can create new types: CREATE TYPE AllRatings as INTEGER
 - Underlying domain is now a new type. Can't compare with INTEGER without a cast. None of the aggregates on INTEGER work on AllRatings

Table Constraints

- More general ICs than key constraints
- Can use queries to express constraint
- Constraints can be named.

```

CREATE TABLE Sailors
( sid INTEGER,
  sname CHAR(10),
  rating INTEGER,
  age REAL,
  PRIMARY KEY (sid),
  CHECK ( rating >= 1
        AND rating <= 10 )

CREATE TABLE Reserves
( sname CHAR(10),
  bid INTEGER,
  day DATE,
  PRIMARY KEY (bid,day),
  CONSTRAINT noInterlakeRes
  CHECK (`Interlake' <>
        ( SELECT B.bname
          FROM Boats B
          WHERE B.bid=bid)))

```

Constraints Over Multiple Relations

- Awkward & Wrong!
- If Sailors is empty, the number of Boats tuples can be anything!
- ASSERTION is the right solution; not associated with either table

```
CREATE TABLE Sailors
```

```
( sid INTEGER,  
  sname CHAR(10),  
  rating INTEGER,  
  age REAL,  
  PRIMARY KEY (sid),  
  CHECK
```

```
((SELECT COUNT (S.sid) FROM Sailors S) +  
 (SELECT COUNT (B.bid) FROM Boats B) < 100)
```

*Number of boats
plus number of
sailors is < 100*

```
CREATE ASSERTION smallClub  
CHECK
```

```
((SELECT COUNT (S.sid) FROM Sailors S) +  
 (SELECT COUNT (B.bid) FROM Boats B) < 100)
```

Triggers

- Trigger: procedure that starts automatically if specified changes occur to the DBMS
- Three parts:
 - Event (activates the trigger)
 - Condition (tests whether the triggers should run)
 - Action (what happens if the trigger runs)
 - Before and After Triggers
- Trigger Execution
 - Row-level Triggers: Once per row
 - Statement-level Triggers: Once per SQL statement

Triggers: Example

```
CREATE TRIGGER init_count BEFORE INSERT ON Students /* Event */
  DECLARE
    count INTEGER /* Action */
  BEGIN
    count := 0
  END
```

```
CREATE TRIGGER incr_count AFTER INSERT ON Student /* Event */
  WHEN (new.age < 18) /* Condition */
  FOR EACH ROW
  BEGIN /* Action */
    count := count + 1;
  END
```

Triggers

- First trigger executed before the activating statement, second executes after the activating statement.
- Options:
 - “BEFORE”
 - “AFTER”
 - “INSTEAD OF” (only valid on views)
- In combination with:
 - “FOR EACH ROW” - execute once per modified record
 - (default) - execute once per activating statement.
Can also specify using “FOR EACH STATEMENT”
- In combination with:
 - “INSERT”
 - “DELETE”
 - “UPDATE”

Triggers

- Referring to values
 - Old
 - New
 - Set of changed record

Triggers: Example

```
CREATE TRIGGER youngSailorUpdate
  AFTER INSERT ON SAILORS
  REFERENCING NEW TABLE NewSailors
  FOR EACH STATEMENT
  INSERT
    INTO YoungSailors(sid, name, age, rating)
  SELECT sid, name, age, rating
  FROM NewSailors N
  WHERE N.age <= 18
```

Triggers v/s Constraints

- Often used to maintain consistency
 - Can you use a foreign key?
 - Foreign keys are not defined operationally
- Constraints are easier to understand than triggers
- Triggers are more powerful.
 - Often used to fill out fields in a form
 - Check complex actions (such as credit limit in a shopping application)
 - Check preferred customer status
 - Generate logs for auditing and security checks.
 - Internally can be used by the DBMS for replication management.