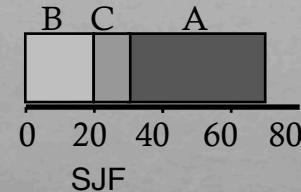
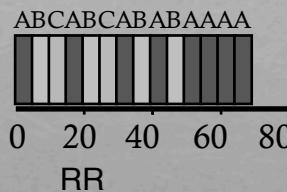


SCHEDULING POLICY: REVIEW

Workload

JOB	arrival	run
A	0	40
B	0	20
C	5	10

Timelines



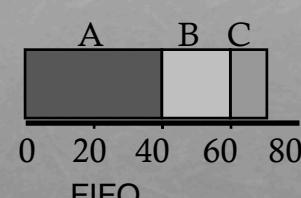
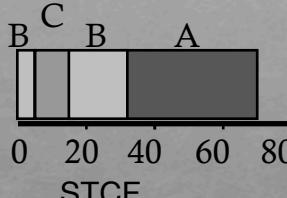
Schedulers:

FIFO

SJF

STCF

RR



UNIVERSITY of WISCONSIN-MADISON
Computer Sciences Department

CS 537
Introduction to Operating Systems

Andrea C. Arpaci-Dusseau
Remzi H. Arpaci-Dusseau

MEMORY VIRTUALIZATION:

Questions answered in this lecture:

What is in the address space of a process (review)?

What are the different ways that that OS can virtualize memory?

Time sharing, static relocation, dynamic relocation

(base, base + bounds, segmentation)

What hardware support is needed for dynamic relocation?

ANNOUNCEMENTS

- Intro to Linux tutorial tonight:
5:30pm in CS 1221
- Switching discussion sections unofficially
- P1
 - Handin directories now available
 - Some test scripts available
 - Due officially by Friday, 5pm; unofficially by Saturday morning, 8am
- Discussion Sections tomorrow
 - P1b questions
 - Watch video ahead of time!
- **Lots of Office Hours**
- [Joe Cai](#)
Office Hours in 5364 CS
Thu 1:20-2:20, Fri 1:20-2:20
- [Ceyhun Alp](#)
Office Hours in 1306 CS
Tue 2:40-3:40, Thu 2:40-3:40
- [Taylor Johnston](#)
Office Hours in 1307 CS
Mon 2:30-3:30, Thu 1:00-2:00
- [Akshay Uttamani](#)
Office Hours in 1302 CS
Wed 2:45-3:45, Fri 3:30-4:30

MORE VIRTUALIZATION

1st part of course: Virtualization

Virtual CPU: *illusion of private CPU registers*

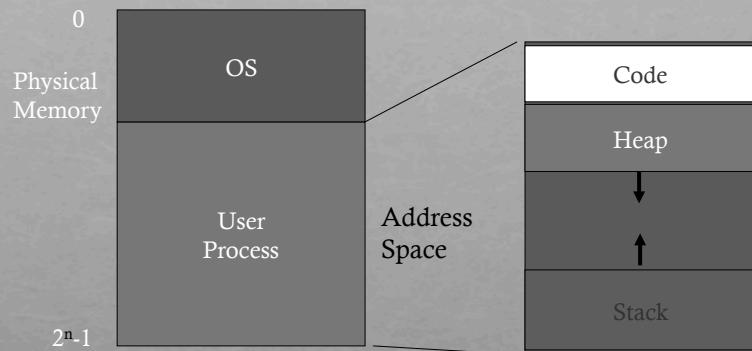
- 2 lectures (mechanism + policy)

Virtual RAM: *illusion of private memory*

- 5 lectures

MOTIVATION FOR VIRTUALIZATION

Uniprogramming: One process runs at a time



Disadvantages:

- Only one process runs at a time
- Process can destroy OS

MULTIPROGRAMMING GOALS

Transparency

- Processes are not aware that memory is shared
- Works regardless of number and/or location of processes

Protection

- Cannot corrupt OS or other processes
- Privacy: Cannot read data of other processes

Efficiency

- Do not waste memory resources (minimize fragmentation)

Sharing

- Cooperating processes can share portions of address space

ABSTRACTION: ADDRESS SPACE

Address space: Each process has set of addresses that map to bytes

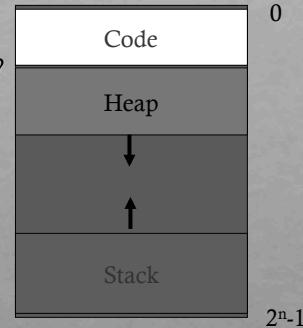
Problem:

How can OS provide illusion of private address space to each process?

Review: What is in an address space?

Address space has static and dynamic components

- Static: Code and some global variables
- Dynamic: Stack and Heap



MOTIVATION FOR DYNAMIC MEMORY

Why do processes need dynamic allocation of memory?

- Do not know amount of memory needed at compile time
- Must be pessimistic when allocate memory statically
 - Allocate enough for worst possible case; Storage is used inefficiently

Recursive procedures

- Do not know how many times procedure will be nested

Complex data structures: lists and trees

- `struct my_t *p = (struct my_t *)malloc(sizeof(struct my_t));`

Two types of dynamic allocation

- Stack
- Heap

STACK ORGANIZATION

Definition: Memory is freed in opposite order from allocation

```
alloc(A);
alloc(B);
alloc(C);
free(C);
alloc(D);
free(D);
free(B);
free(A);
```

Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation

WHERE ARE STACKS USED?

OS uses stack for procedure call frames (local variables and parameters)

```
main () {
    int A = 0;
    foo (A);
    printf("A: %d\n", A);
}

void foo (int Z) {
    int A = 2;
    Z = 5;
    printf("A: %d Z: %d\n", A, Z);
}
```

HEAP ORGANIZATION

Definition: Allocate from any random location: malloc(), new()

- Heap memory consists of allocated areas and free areas (holes)
- Order of allocation and free is unpredictable

Advantage

- Works for all data structures

Disadvantages

- Allocation can be slow
- End up with small chunks of free space - fragmentation
- Where to allocate 12 bytes? 16 bytes? 24 bytes??
- What is OS's role in managing heap?
- OS gives big chunk of free memory to process; library manages individual allocations



QUIZ: MATCH THAT ADDRESS LOCATION

```
int x;
int main(int argc, char *argv[]) {
    int y;
    int *z = malloc(sizeof(int));;
}
```

Possible segments: static data, code, stack, heap

What if no static data segment?

Address	Location
x	Static data → Code
main	Code
y	Stack
z	Stack
*z	Heap

MEMORY ACCESSES

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    int x;
    x = x + 3;
}
```

`otool -tv demo1.o
(or objdump on Linux)`



```
0x10:  movl 0x8(%rbp), %edi
0x13:  addl $0x3, %edi
0x19:  movl %edi, 0x8(%rbp)
```

%rbp is the base pointer:
points to base of current stack frame

QUIZ: MEMORY ACCESSES?

Initial %rip = 0x10
%rbp = 0x200

→ 0x10: movl 0x8(%rbp), %edi
0x13: addl \$0x3, %edi
0x19: movl %edi, 0x8(%rbp)

%rbp is the base pointer:
points to base of current stack frame

%rip is instruction pointer (or program counter)

Fetch instruction at addr 0x10
Exec:

load from addr 0x208

Fetch instruction at addr 0x13
Exec:

no memory access

Fetch instruction at addr 0x19
Exec:

store to addr 0x208

Memory Accesses to what addresses?

HOW TO VIRTUALIZE MEMORY?

Problem: How to run multiple processes simultaneously?

Addresses are “hardcoded” into process binaries

How to avoid collisions?

Possible Solutions for Mechanisms (covered today):

1. Time Sharing
2. Static Relocation
3. Base
4. Base+Bounds
5. Segmentation

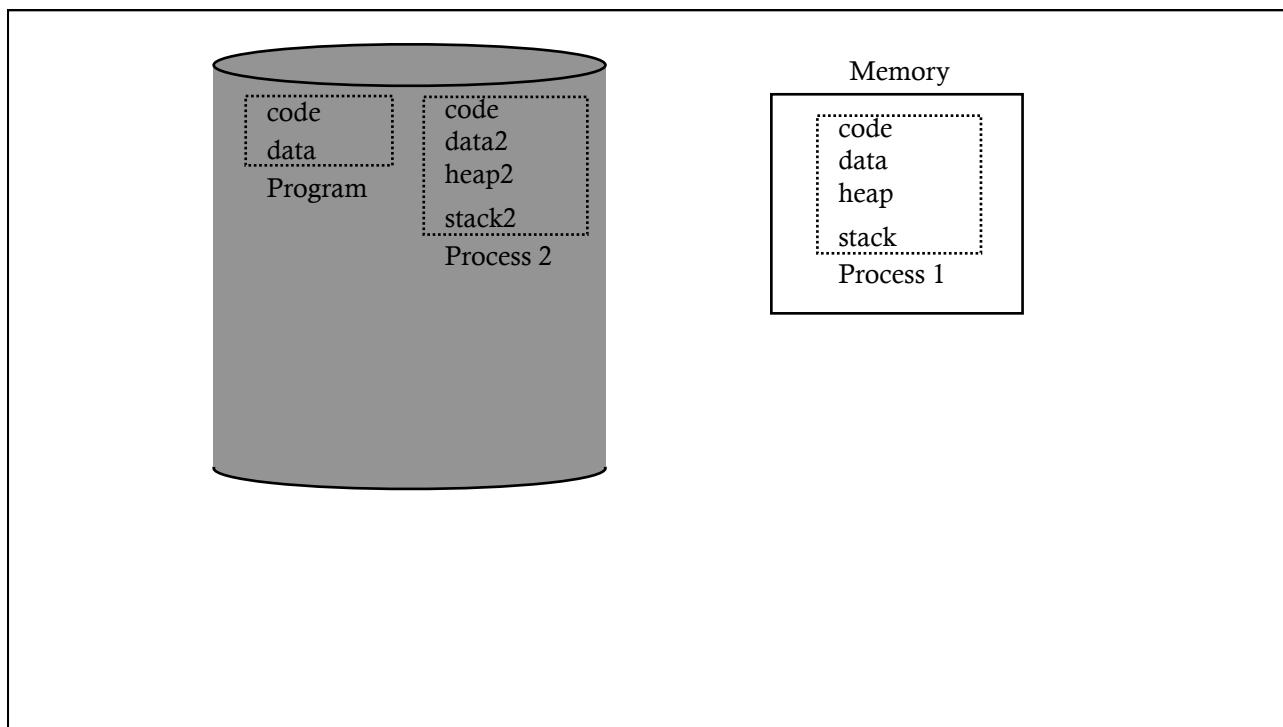
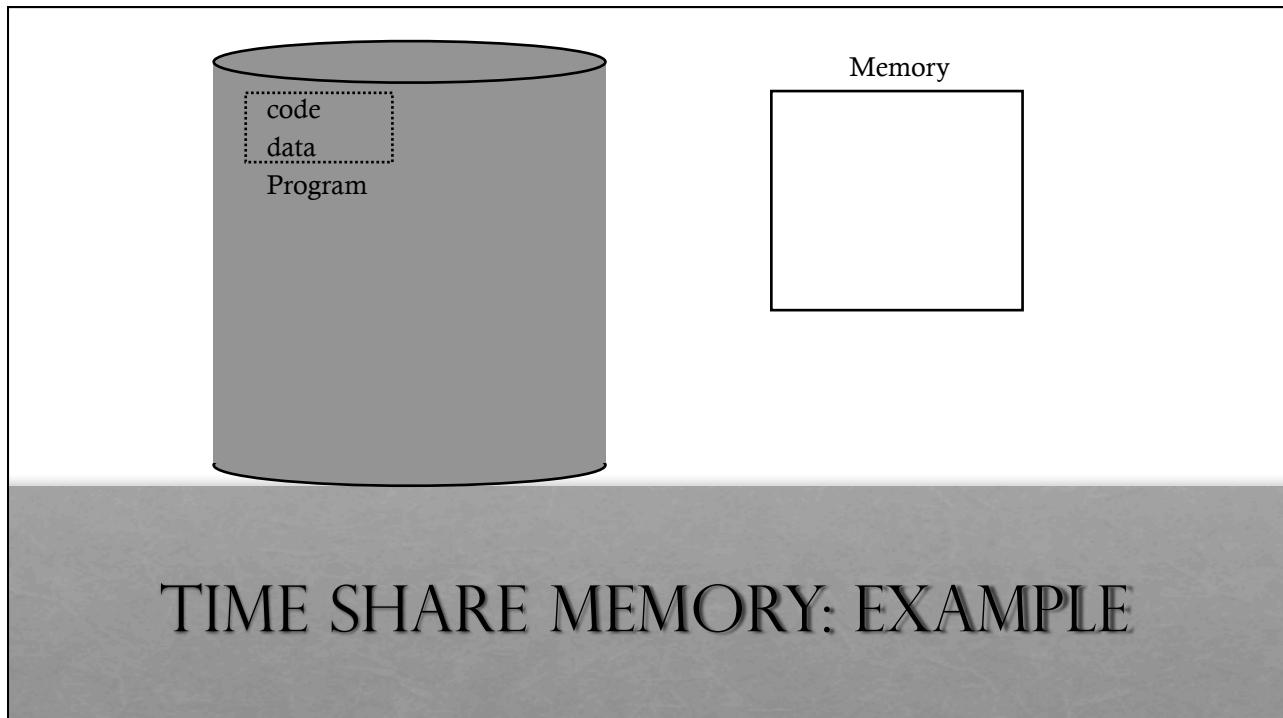
1) TIME SHARING OF MEMORY

Try similar approach to how OS virtualizes CPU

Observation:

OS gives illusion of many virtual CPUs by saving **CPU registers** to **memory** when a process isn't running

Could give illusion of many virtual memories by saving **memory** to **disk** when process isn't running



PROBLEMS WITH TIME SHARING MEMORY

Problem: Ridiculously poor performance

Better Alternative: space sharing

- At same time, space of memory is divided across processes

Remainder of solutions all use space sharing

2) STATIC RELOCATION

- Idea: OS rewrites each program before loading it as a process in memory
- Each rewrite for different process uses different addresses and pointers
- Change jumps, loads of static data

```

        0x1010: movl 0x8(%rbp), %edi
        0x1013: addl $0x3, %edi
        0x1019: movl %edi, 0x8(%rbp)

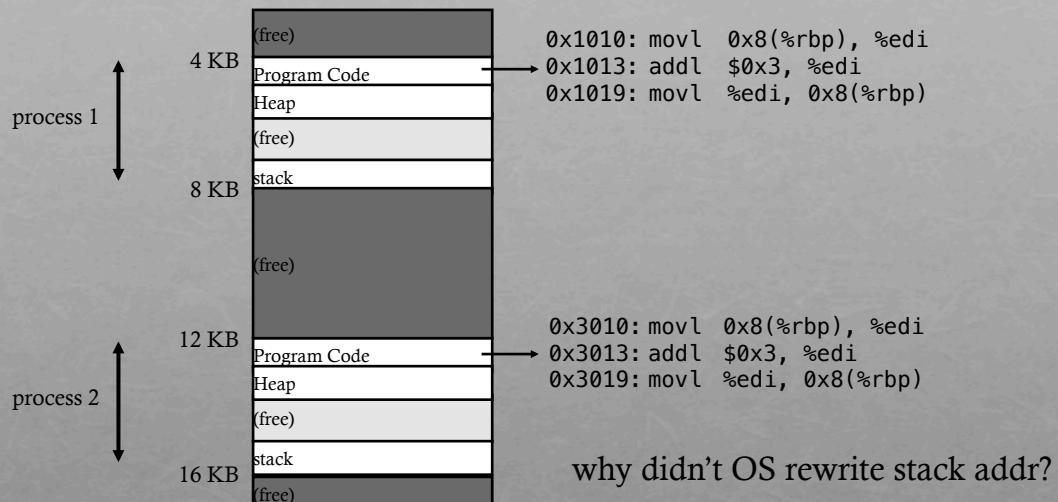
    rewrite →

    • 0x10: movl 0x8(%rbp), %edi
    • 0x13: addl $0x3, %edi
    • 0x19: movl %edi, 0x8(%rbp)

        0x3010: movl 0x8(%rbp), %edi
        0x3013: addl $0x3, %edi
        0x3019: movl %edi, 0x8(%rbp)

    rewrite →
  
```

STATIC: LAYOUT IN MEMORY



STATIC RELOCATION: DISADVANTAGES

No protection

- Process can destroy OS or other processes
- No privacy

Cannot move address space after it has been placed

- May not be able to allocate new process

3) DYNAMIC RELOCATION

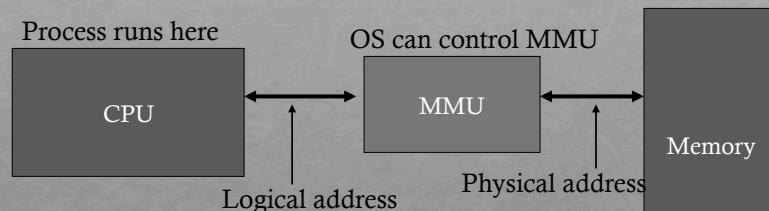
Goal: Protect processes from one another

Requires hardware support

- Memory Management Unit (MMU)

MMU dynamically changes process address at every memory reference

- Process generates logical or virtual addresses (in their address space)
- Memory hardware uses physical or real addresses



HARDWARE SUPPORT FOR DYNAMIC RELOCATION

Two operating modes

- Privileged (protected, kernel) mode: OS runs
 - When enter OS (trap, system calls, interrupts, exceptions)
 - Allows certain instructions to be executed
 - Can manipulate contents of MMU
 - Allows OS to access all of physical memory
- User mode: User processes run
 - Perform translation of logical address to physical address

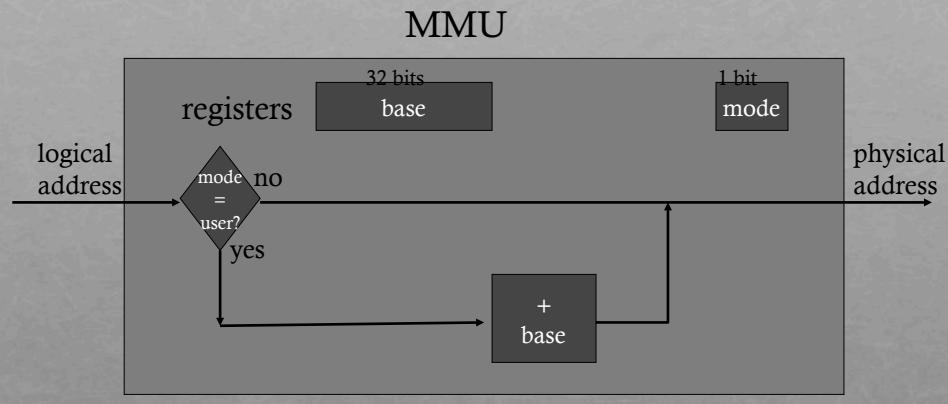
Minimal MMU contains **base register** for translation

- base: start location for address space

IMPLEMENTATION OF DYNAMIC RELOCATION: BASE REG

Translation on every memory access of user process

- MMU adds base register to logical address to form physical address

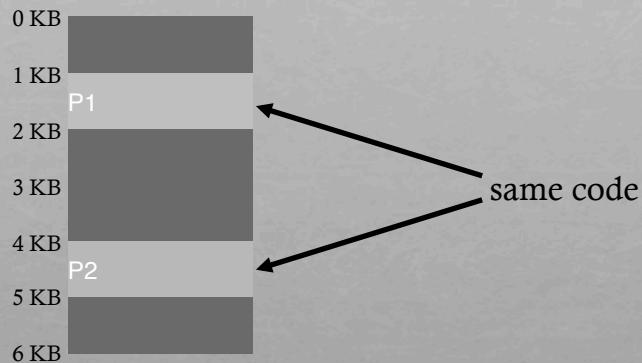


DYNAMIC RELOCATION WITH BASE REGISTER

Idea: translate virtual addresses to physical by adding a fixed offset each time.

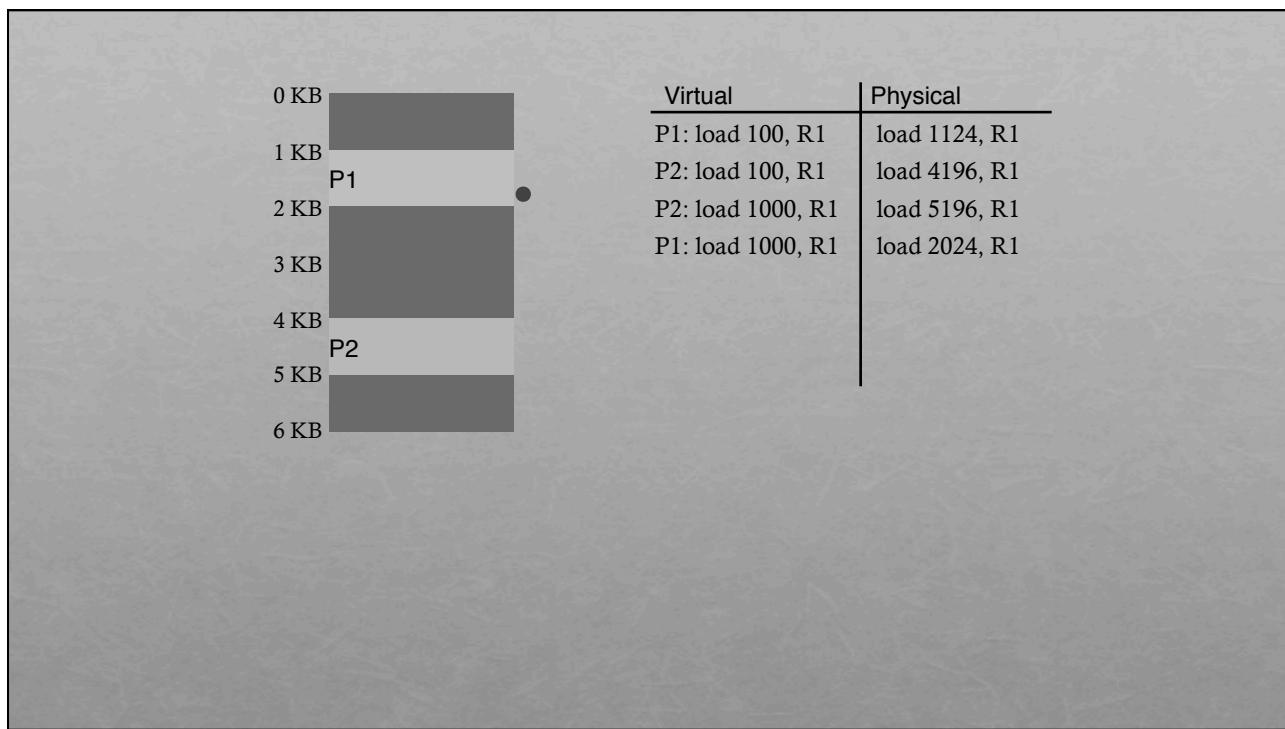
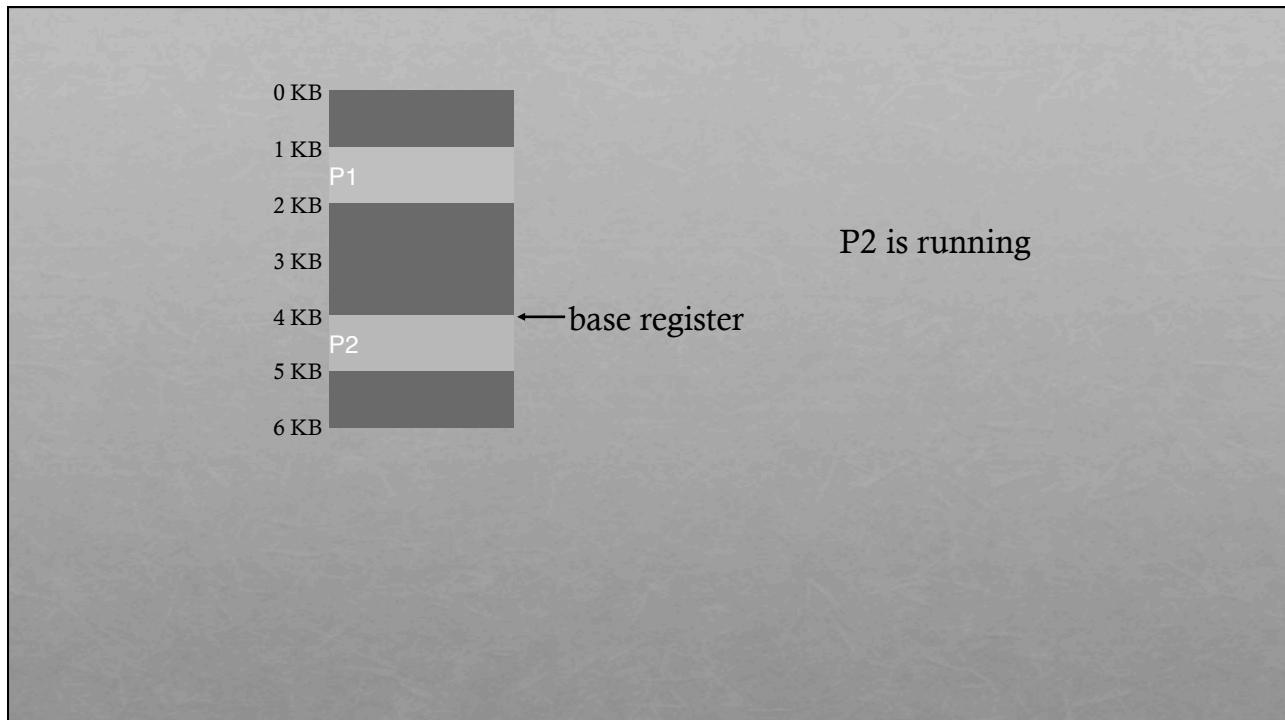
Store offset in base register

Each process has different value in base register



VISUAL EXAMPLE OF DYNAMIC RELOCATION: BASE REGISTER





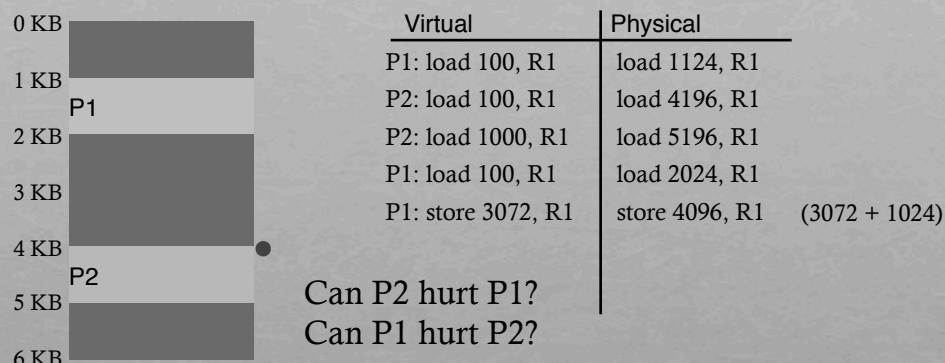
QUIZ: WHO CONTROLS THE BASE REGISTER?

What entity should do translation of addresses with base register?

- (1) process, (2) OS, or (3) HW

What entity should modify the base register?

- (1) process, (2) OS, or (3) HW



How well does dynamic relocation do with base register for protection?

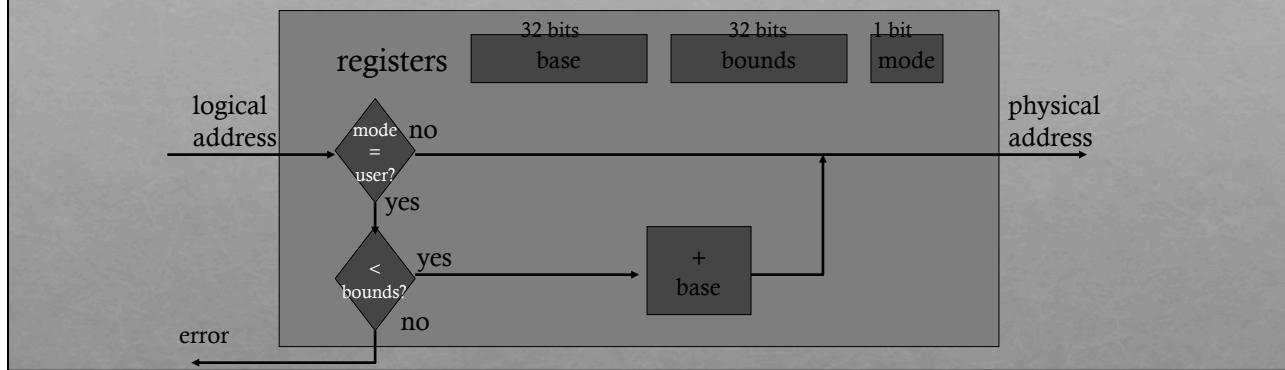
4) DYNAMIC WITH BASE+BOUNDS

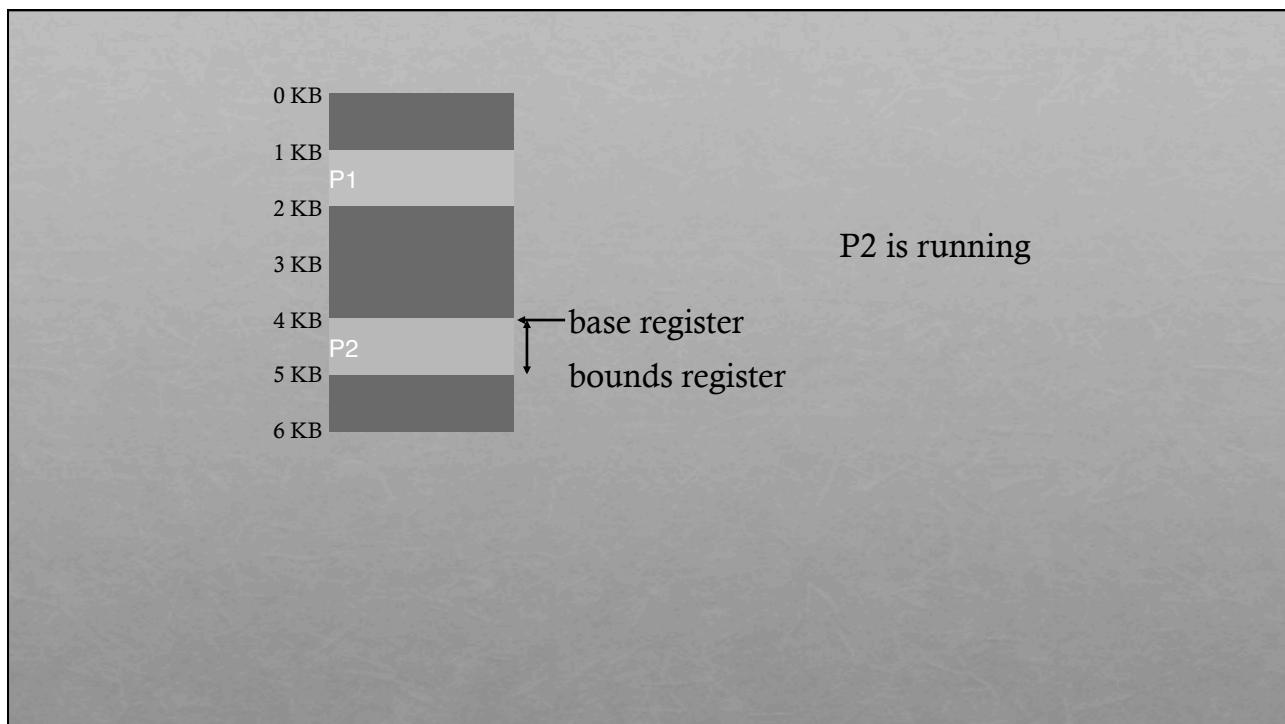
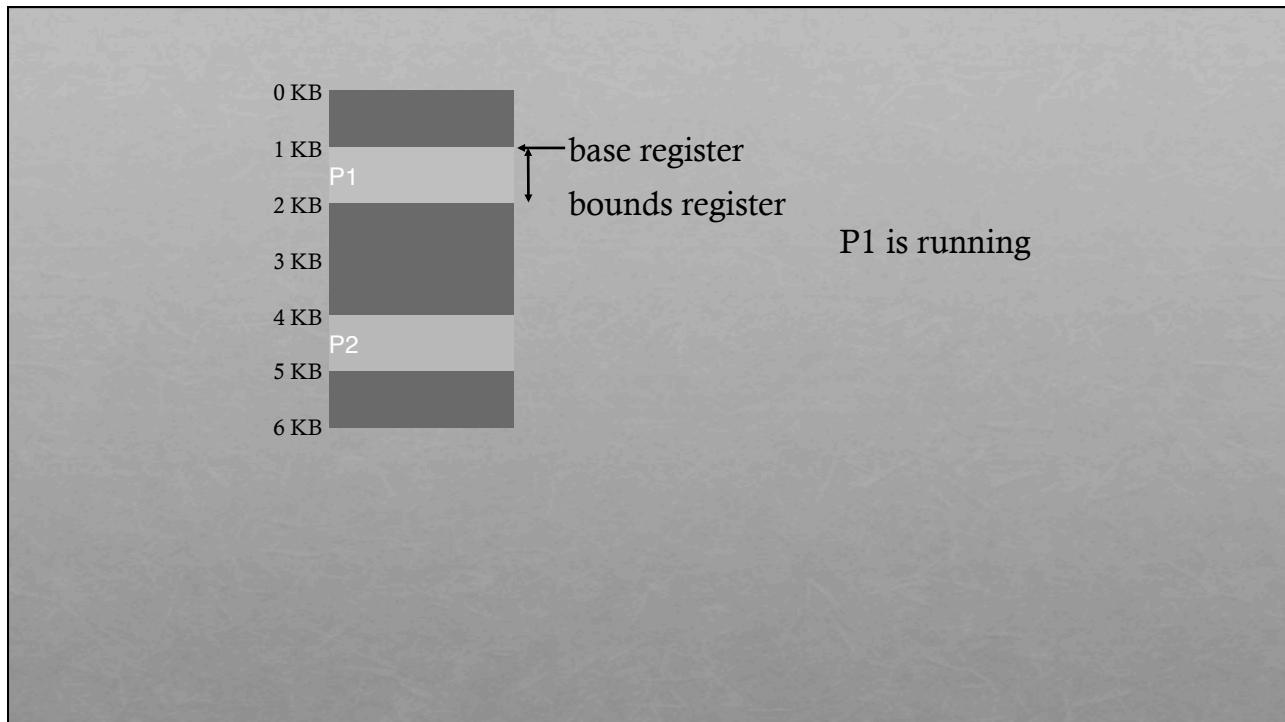
- Idea: limit the address space with a bounds register
- Base register: smallest physical addr (or starting location)
- Bounds register: size of this process's virtual address space
 - Sometimes defined as largest physical address (base + size)
- OS kills process if process loads/stores beyond bounds

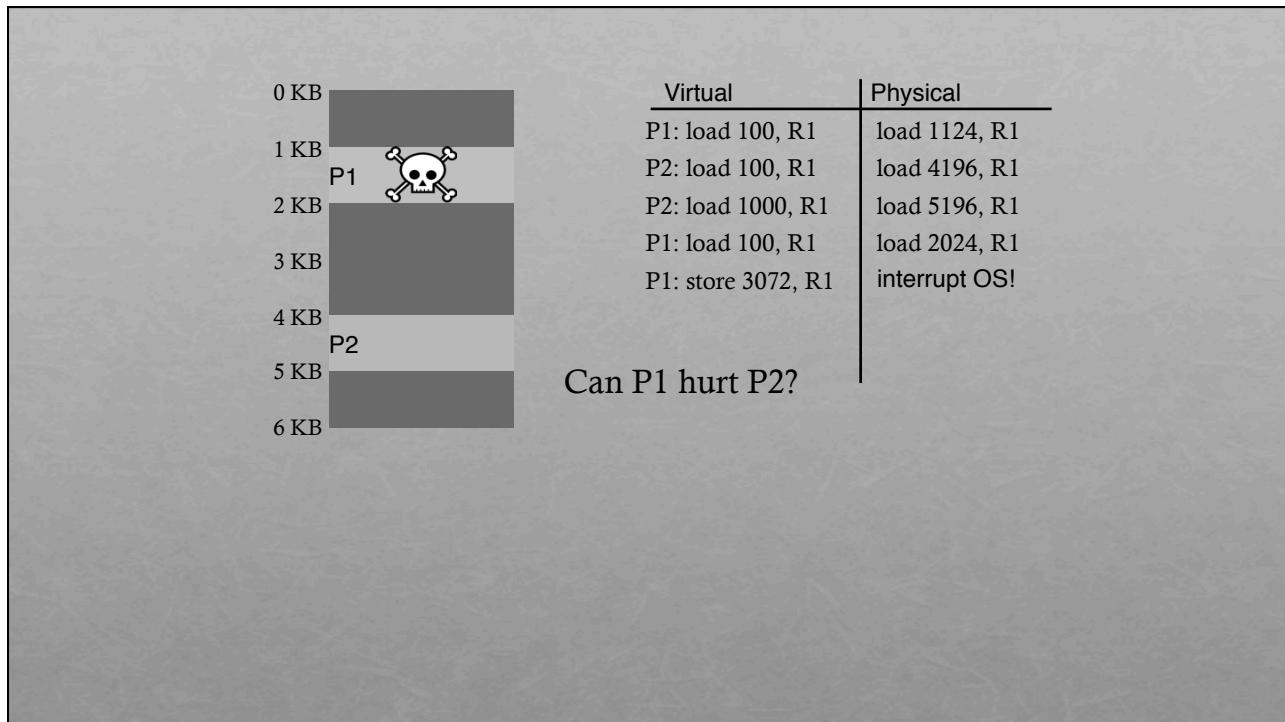
IMPLEMENTATION OF BASE+BOUNDS

Translation on every memory access of user process

- MMU compares logical address to bounds register
 - if logical address is greater, then generate error
- MMU adds base register to logical address to form physical address







MANAGING PROCESSES WITH BASE AND BOUNDS

Context-switch

- Add base and bounds registers to PCB
- Steps
 - Change to privileged mode
 - Save base and bounds registers of old process
 - Load base and bounds registers of new process
 - Change to user mode and jump to new process

What if don't change base and bounds registers when switch?

Protection requirement

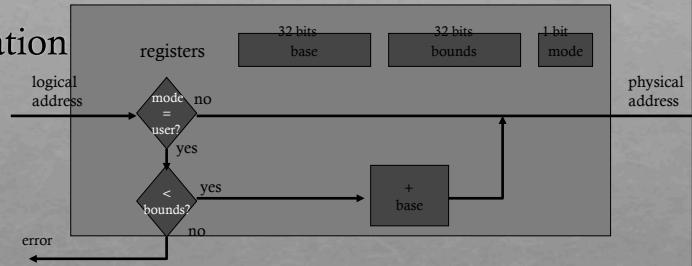
- User process cannot change base and bounds registers
- User process cannot change to privileged mode

BASE AND BOUNDS ADVANTAGES

Advantages

- Provides protection (both read and write) across address spaces
- Supports dynamic relocation
 - Can place process at different locations initially and also move address spaces

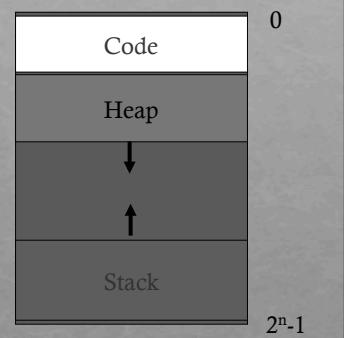
- Simple, inexpensive implementation
 - Few registers, little logic in MMU
- Fast
 - Add and compare in parallel



BASE AND BOUNDS DISADVANTAGES

Disadvantages

- Each process must be allocated contiguously in physical memory
 - Must allocate memory that may not be used by process
- No partial sharing: Cannot share limited parts of address space



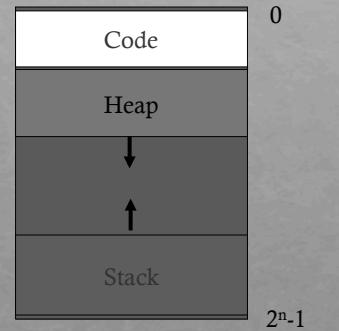
5) SEGMENTATION

Divide address space into logical segments

- Each segment corresponds to logical entity in address space
 - code, stack, heap

Each segment can independently:

- be placed separately in physical memory
- grow and shrink
- be protected (separate read/write/execute protection bits)



SEGMENTED ADDRESSING

Process now specifies segment and offset within segment

How does process designate a particular segment?

- Use part of logical address
 - Top bits of logical address select segment
 - Low bits of logical address select offset within segment

What if small address space, not enough bits?

- Implicitly by type of memory reference
- Special registers

SEGMENTATION IMPLEMENTATION

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R W
0	0x2000	0x6ff	1 0
1	0x0000	0x4ff	1 1
2	0x3000	0xffff	1 1
3	0x0000	0x000	0 0

remember:
1 hex digit->4 bits

QUIZ: ADDRESS TRANSLATIONS WITH SEGMENTATION

MMU contains Segment Table (per process)

- Each segment has own base and bounds, protection bits
- Example: 14 bit logical address, 4 segments; how many bits for segment? How many bits for offset?

Segment	Base	Bounds	R W
0	0x2000	0x6ff	1 0
1	0x0000	0x4ff	1 1
2	0x3000	0xffff	1 1
3	0x0000	0x000	0 0

remember:
1 hex digit->4 bits

Translate logical addresses (in hex) to physical addresses

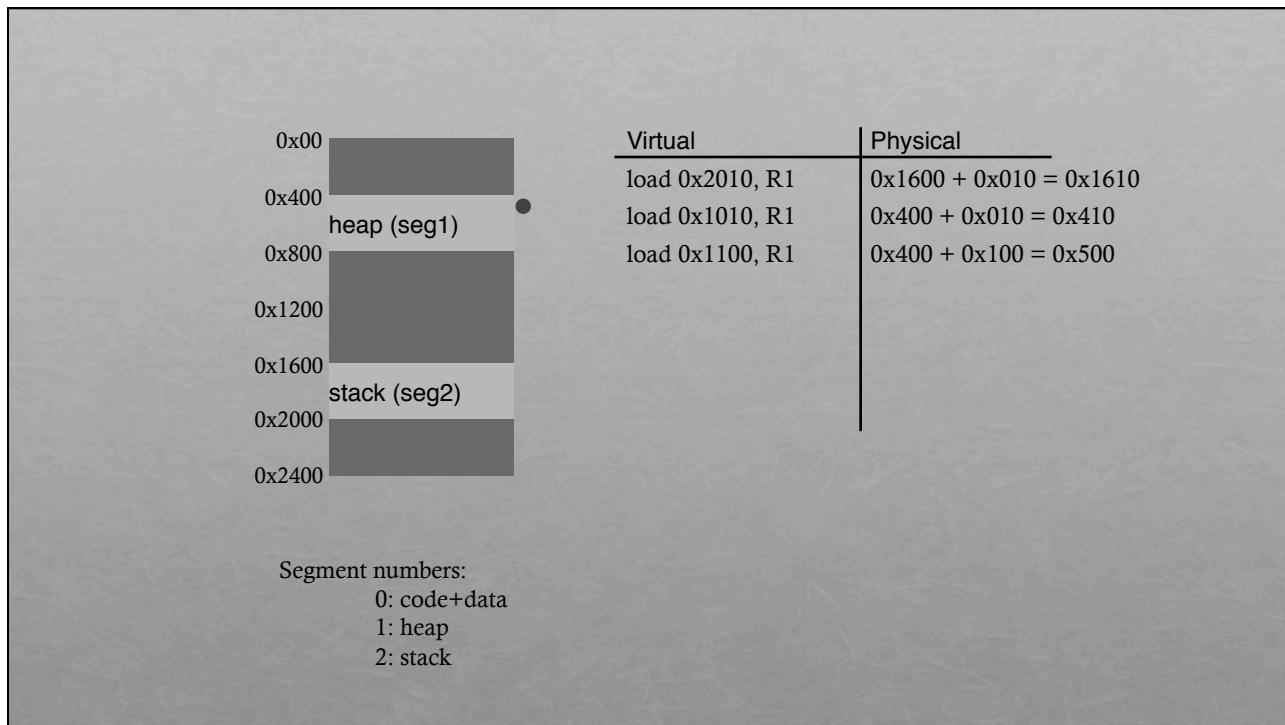
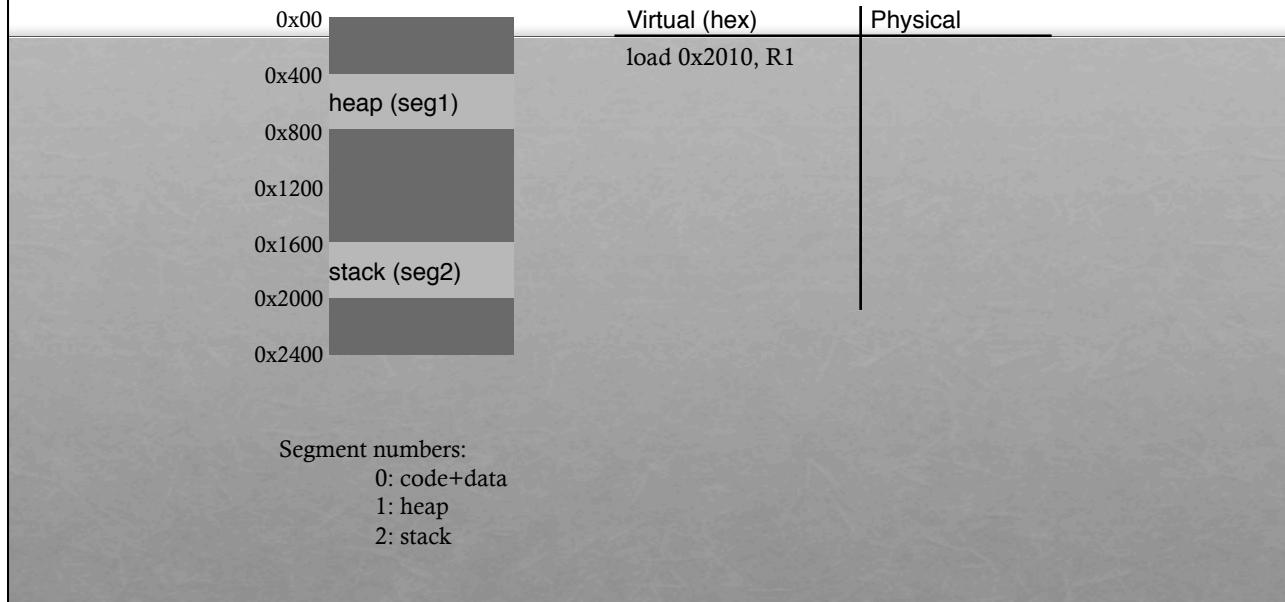
0x0240:

0x1108:

0x265c:

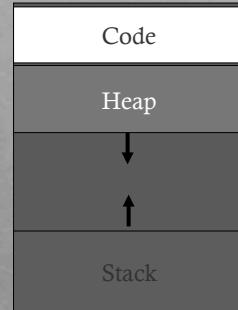
0x3002:

VISUAL INTERPRETATION



ADVANTAGES OF SEGMENTATION

- Enables sparse allocation of address space
 - Stack and heap can grow independently
 - Heap: If no data on free list, dynamic memory allocator requests more from OS (e.g., UNIX: malloc calls sbrk())
 - Stack: OS recognizes reference outside legal segment, extends stack implicitly
- Different protection for different segments
 - Read-only status for code
- Enables sharing of selected segments
- Supports dynamic relocation of each segment



DISADVANTAGES OF SEGMENTATION

Each segment must be allocated contiguously

- May not have sufficient physical memory for large segments

Fix in next lecture with paging...

CONCLUSION

HW+OS work together to virtualize memory

- Give illusion of private address space to each process

Add MMU registers for base+bounds so translation is fast

- OS not involved with every address translation, only on context switch or errors

Dynamic relocation with segments is good building block

- Next lecture: Solve fragmentation with paging